



11.6.2018

München / Microservice Summit

Kubernetes - the abstract cloud

Jörg Müller - @joergm

INNOQ

- **architecture,
development,
devOps**
- **focus on
platform &
infrastructure**



Jörg Müller

Principal Consultant
innoQ Deutschland GmbH

joerg.mueller@innoq.com
@joergm

What to expect?

- Overview, Ways to get Kubernetes and basic concepts
- Core abstractions
- Internal Architecture
- Deploying complex applications
- Production readiness

Timeslots

- 9:30 - 10:30 Slot 1
- 10:30 - 11:00 Coffee break
- 11:00 - 12:30 Slot 2
- 12:30 - 13:30 Lunch
- 13:30 - 15:00 Slot 3
- 15:00 - 15:30 Coffee break
- 15:30 - 17:00 Slot 4

Prerequisites & rules

- Kubernetes know-how not necessary, but it doesn't hurt
- Basic knowledge about Docker is assumed
- Demos can be followed but don't have to
 - github.com/JoergM/kubernetes_workshop_demos
- Please ask questions!



Kubernetes Overview

Docker Recap

Docker container at runtime

- Isolated process
- Separate file system
- Own network address and port space

Docker container - advantages

- **Better isolation than package management on same machine**
 - **e.g. multiple versions of core libraries**
 - **not necessary to coordinate available ports**
- **Faster startup than virtual machine images**
- **Better resource usage compared to VMs**

Docker images

- Standardized format
- Container hierarchies and difference file system
- Registries
- Unique name format
(Registry/username/imagename:version)
- Simple Text-Format to create new images (Dockerfile)

Docker images - advantages

- **Deployment format independent of implementation technology**
- **Same deliverable in all stages (Development, CI, Tests, Production)**
- **Container hierarchies allow simpler patch management**
- **Definition simpler than most package manager definitions**

What is Kubernetes?

Kubernetes — adds to Docker

- Handling of multiple servers
 - Scheduling of containers
 - Networking
 - Failure handling
- Service Discovery features
- Many other useful abstractions for container interactions

Kubernetes - executive summary

- **Kubernetes (K8s) is an open-source system for automating deployment, scaling, and management of containerized applications**
- **Marketing claim:**
 - **Planet Scale**
 - **Never Outgrow**
 - **Run Anywhere**

Kubernetes — brief history

- Designed by Google, later donated to Cloud Native Computing Foundation
- Heavily influenced by Google's internal Borg system
- Code name: Project Seven
- Initial release: 7 June 2014 / 15 December 2015 (first stable version)

Why abstract cloud?

Why do we need one?

- Working on local machines
- Prevent Vendor lock in
 - Less specific Know How necessary
 - Easier to move
- Common way to automate complex setups
 - For inhouse applications
 - Also for software vendors

Kubernetes

- You define resources needed not machines or implementations
- Kubernetes manages resources
- Has a large base of runtime environments

Application

Resource abstraction

Resource management

Runtime environment

We tried that before ...

- **Virtual machines**
- **Configuration management (Puppet, Ansible, Chef)**
- **Terraform, CloudFormation**
- **PaaS**
- **...**

Standing on Shoulders

- Container abstractions
- Googles experiences running Borg
- Focus on immutable infrastructure

K8s for microservices

Challenges

- **Deployment**
- **Configuration**
- **Service Discovery**
- **Load Balancing**
- **Routing**
- **Resilience**

Kubernetes

- Provides solutions for those challenges
- Is available everywhere
- Becomes more and more widespread
 - So developers know how to solve those challenges
 - Operations accepts and knows the solution
- Microservices infrastructure loses a lot of its horror



Ways to get Kubernetes

Local installation

Installing a simple Kubernetes on your notebook.

- Minikube
- Docker native

Online Tryout

Try Kubernetes without installing anything.

- <https://www.katacoda.com/courses/kubernetes/playground>
- <https://labs.play-with-k8s.com/>

By Cloud providers

Managed Kubernetes is now offered by all major cloud providers.

- **Google Kubernetes Engine (GKE)**
- **Azure AKS**
- **IBM Cloud Kubernetes Services**
- **Amazon EKS (GA just started in us-east and us-west)**
- **Digital Ocean Kubernetes (coming soon)**

Specialised Kubernetes providers

Offering managed Kubernetes on different platforms.
Often including Support and On-Premise install.

- GiantSwarm
- Rancher
- Tectonic
- Kontena Pharos
- ...

PaaS Solutions

Platform as a Service built on Kubernetes or offering Kubernetes services.

- **RedHat OpenShift**
- **CloudFoundry Container runtime**

Self install

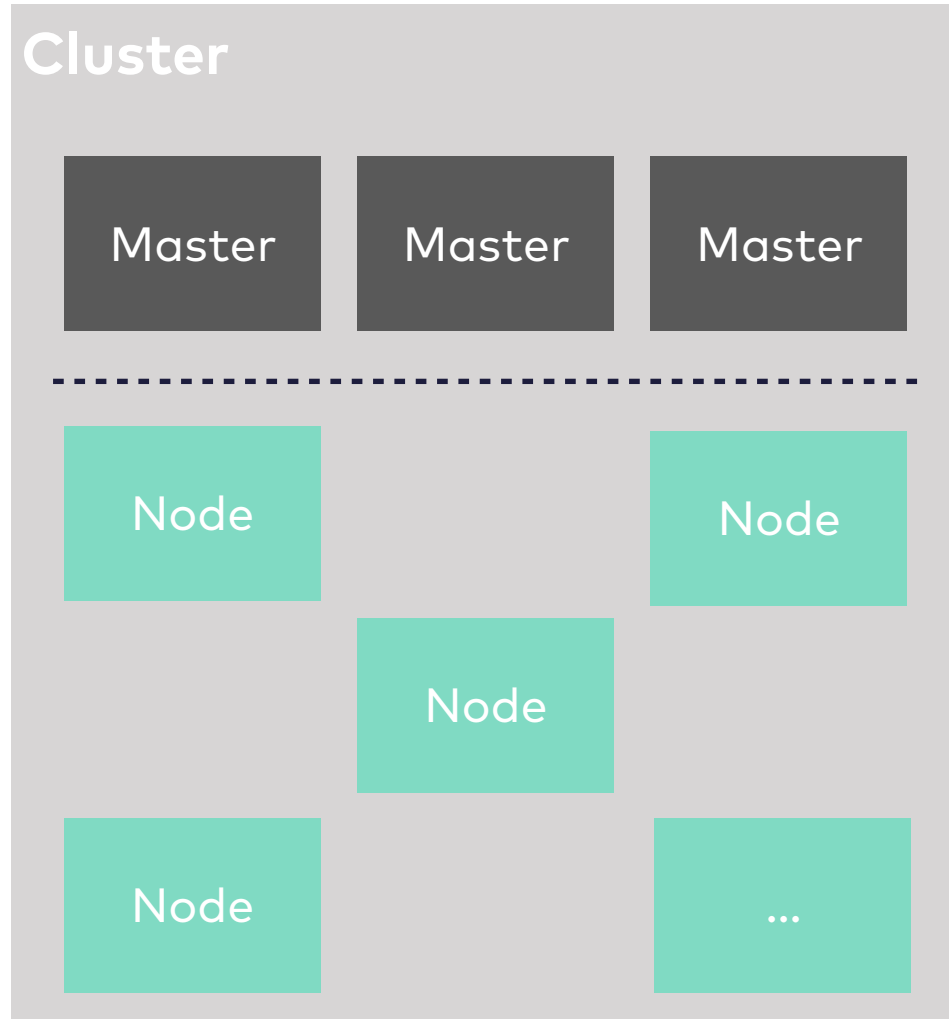
Finally a lot of options to install Kubernetes yourself on Cloud Providers or On-Premise.

- kubeadm
- KOPS
- <https://github.com/kelseyhightower/kubernetes-the-hard-way>




Basic concepts

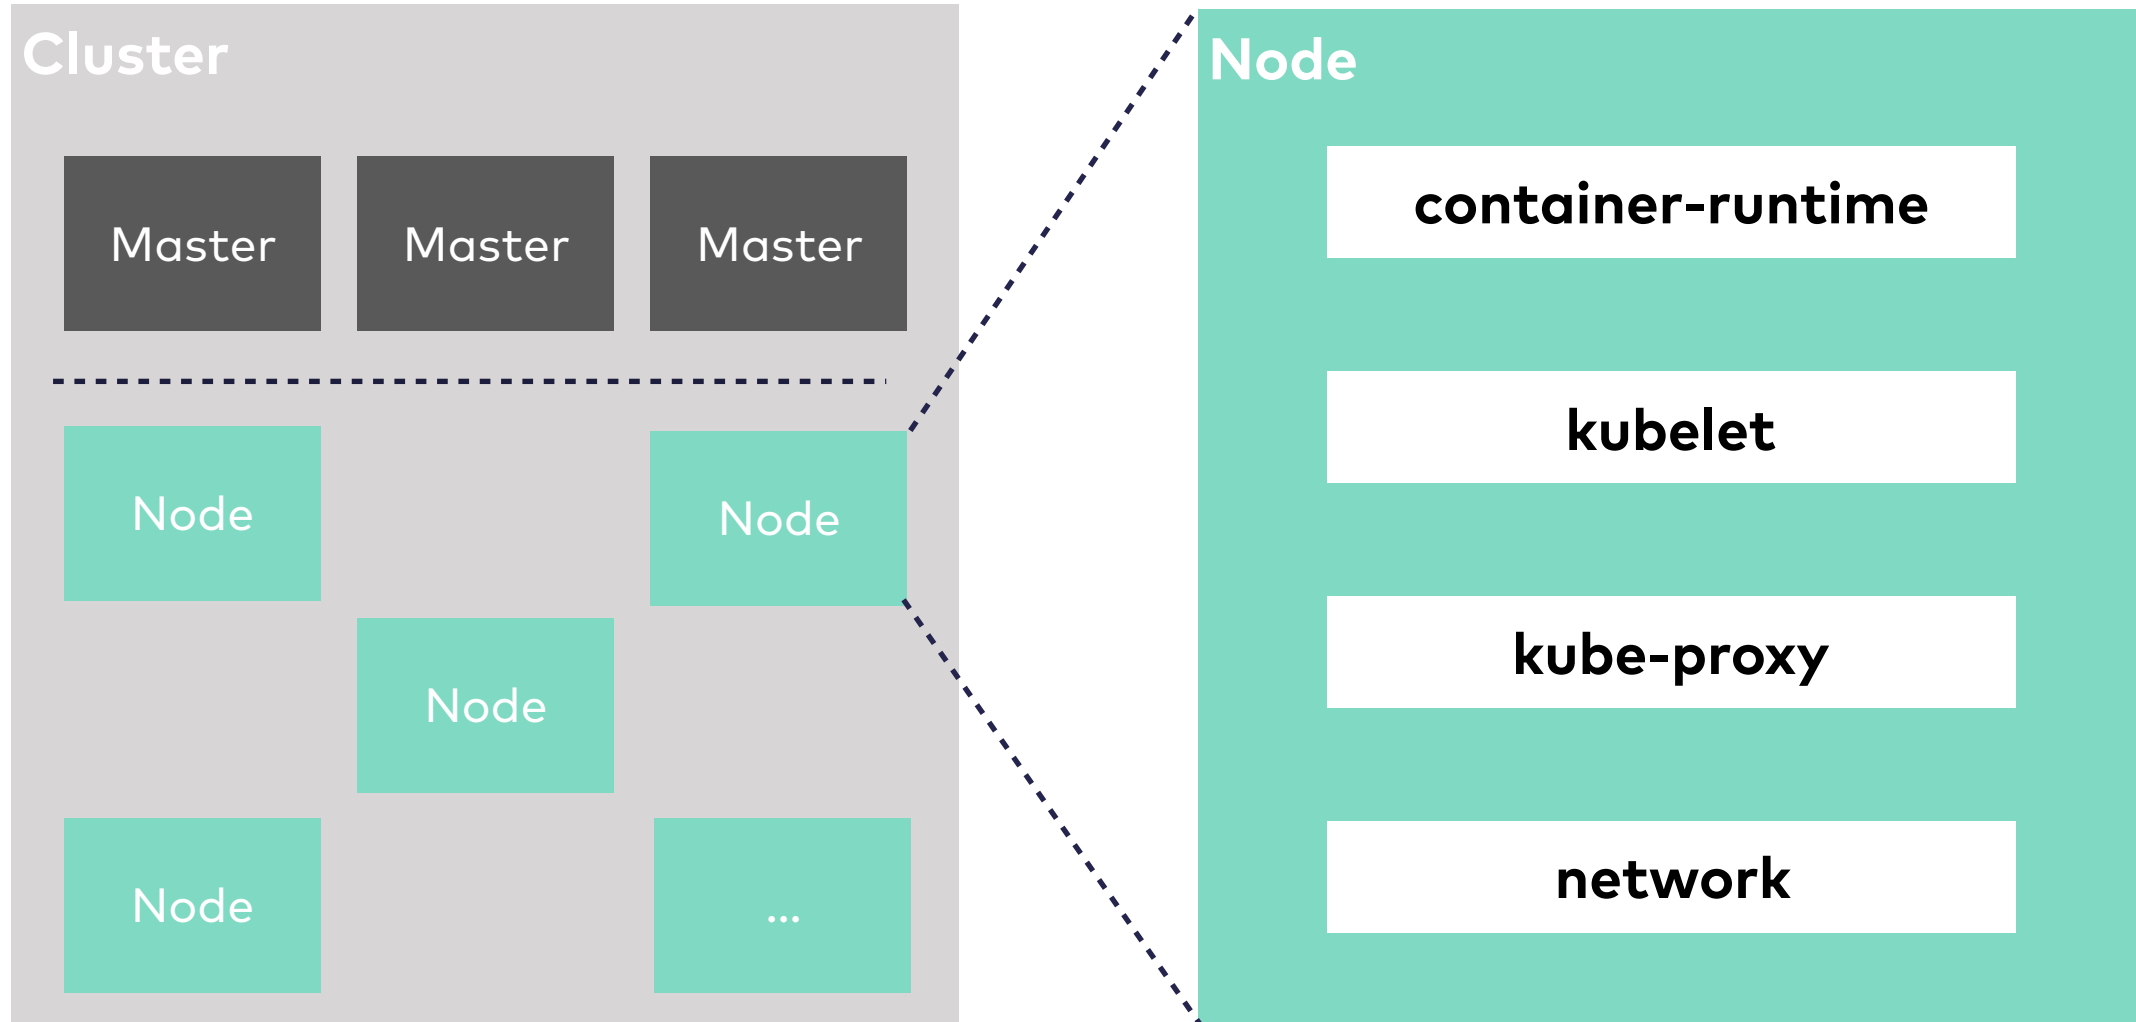
Cluster overview



Master Components



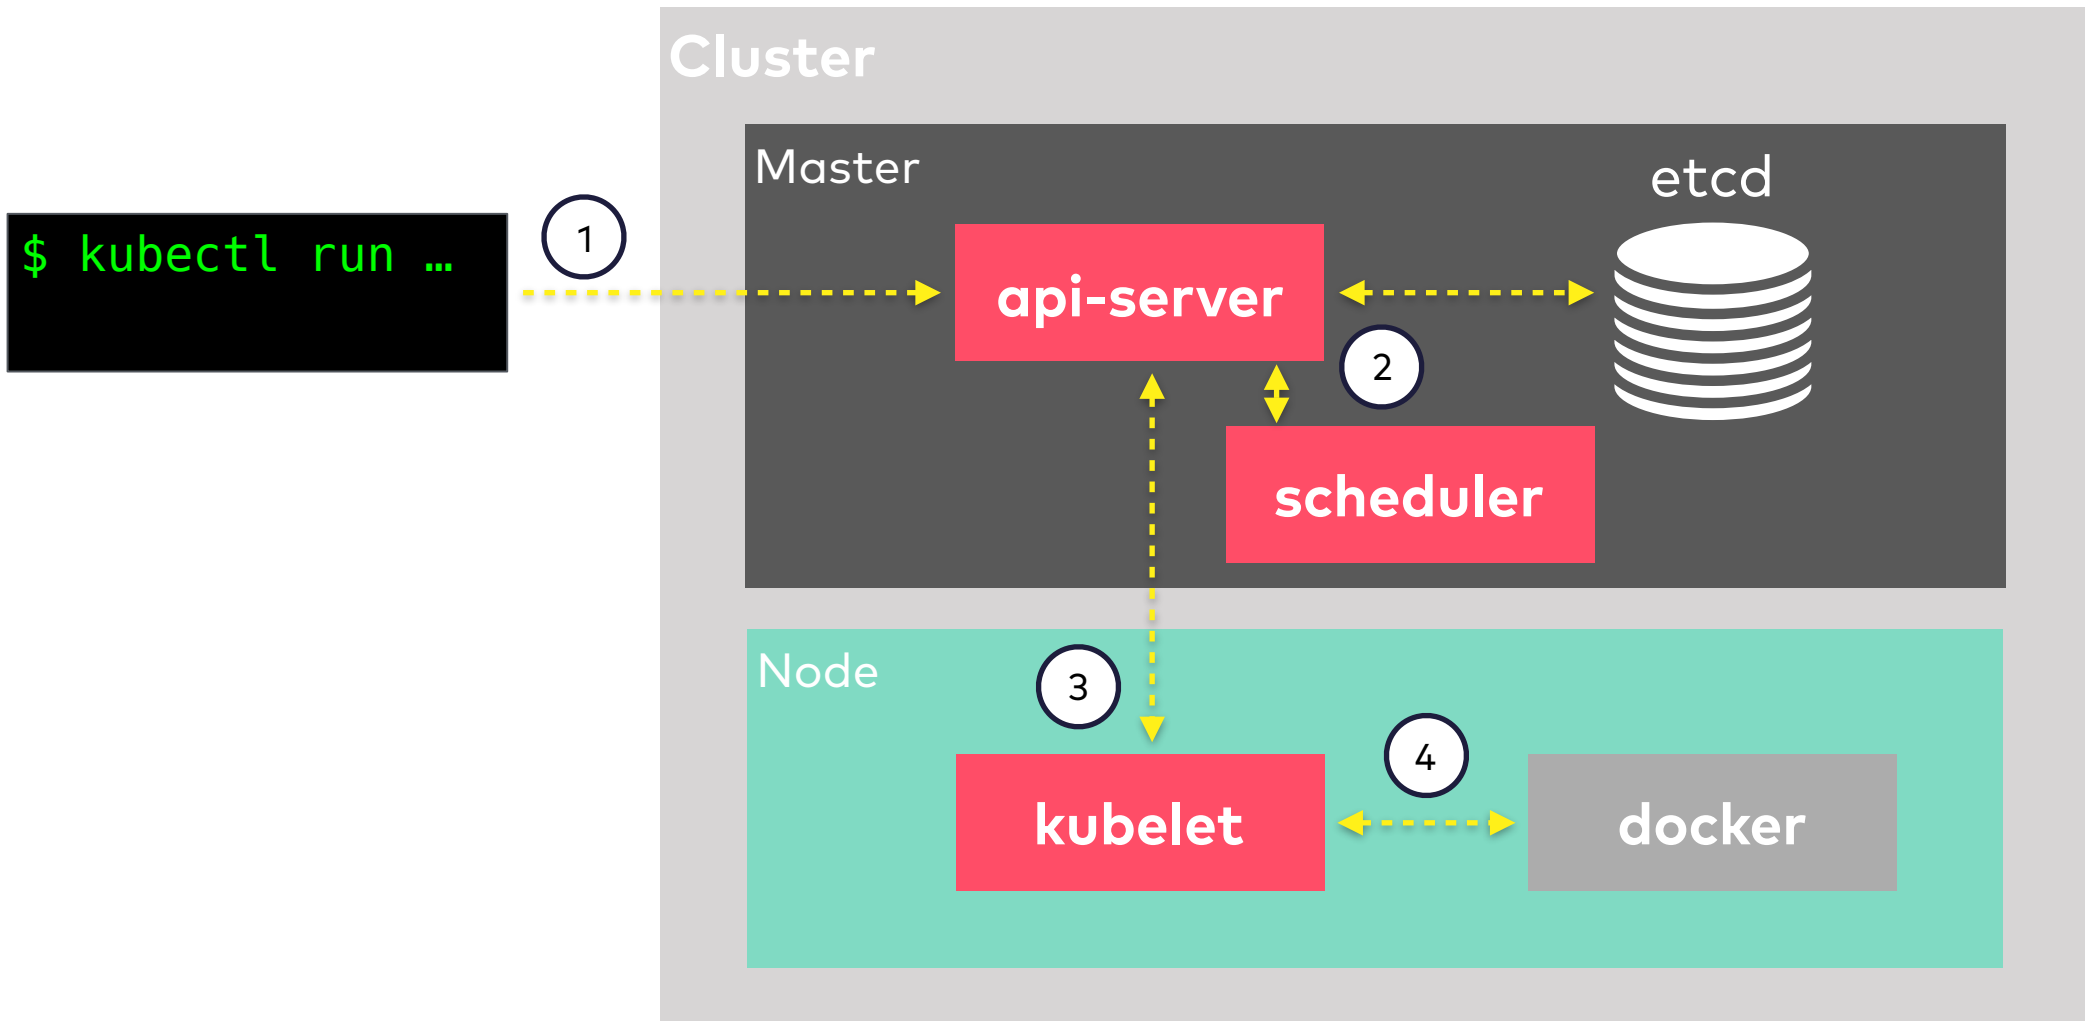
Node Components



API Objects

- Persistent (in etcd)
- represent the **desired state** of the cluster
- Have
 - Spec
 - Status

API Objects interaction



#Computing #Software #Interface #BigData
#ComputerGeneratedDesign

*1981 in London (Glt.
lebt und arbeitet in London

-1982 in London (20).
lives and works in London

085

2011 fortlaufend, Anwendung, 3D-Druck-
Flachfiguren, 3D-Drucke, Tablet; mit
Ständern, Monitore
Webseite: tomeaweb.com

2017 ongoing: Web application, 3-D printed
plastic figures, 3-D prints, tablets on stands,
monitors
Website: www.wmccr.com

Vertigo in the Face of the Infinite fast verschiedene schöpferische und miteinander in Wechselwirkung stehende digitale und materielle Komponenten zusammen. Die Suite im Internet verbindet 3-D-animierte Bewegtbilder, vornehmlich hoch auf und man kann endlich an der Anfangsposition. Über die Kamera führt eigenen Gedanken können die Besucher immer die Website frei Genießen dem digitalen Saal hinzufügen. Die Website findet in der Ausstellung seine physische Entsprechung in dem beeindruckenden Tausend für den Wasserwurf und einer Reihe von synchronisierten 3-D-Druckern, die eine enorme Dimension schaffen das

maschinellen Prozessen und den Maschinen, die über den Herstellungsprozess von On-demand-Artikeln und einer virtuellen Endloskette in 3-D miteinander verflochten sind. Die Arbeit greift die umfassende menschliche Interaktion mit den unterbrochenen Armen, die das Internet und mit den sozialen Medien verbindet.

Vertigo in the Face of the Infinite describes various interrelated and productive digital and physical elements. A stack of found objects is arranged vertically in height, comprised of S-D (Solid-Digital) modules, and is made of wood and veneers. It is an infinite spiral staircase, where the steps are not always even and can be used to sit, but can have their own captured and stored in the lower end of the columns on their surfaces. The staircase also has a physical element in the exhibition, comprised of glass-cast objects and a showcase of S-D prints, from the drawing website www.solid-digital.com as well as the drawing website as a full, variable multi-screen display.

The network approach to the relationships between operating procedures and the objects involved in the production process, associated in the production of objects made in the factory, is a healthy infinite 3-D column. The network receives the wider human engagement with adopting web-based newfeeds and social media where information is vertically stacked and open for contributions in order to stimulate continued use. The attitudes that such newfeeds achieve will create a new kind of vertigo of information.



**github.com/JoergM/
kubernetes_workshop_demos**

Pods


Pod

- **Deployment-Unit in Kubernetes**
- **A pod consists of one or more containers**
- **Containers in a pod share network**
- **Containers in a pod can share volumes**
- **Each pod receives its own cluster-wide and cluster internal IP address**


Pod with a single container



Sharing network




Sharing volumes



Pods with init containers



Complex pod patterns



DEMO

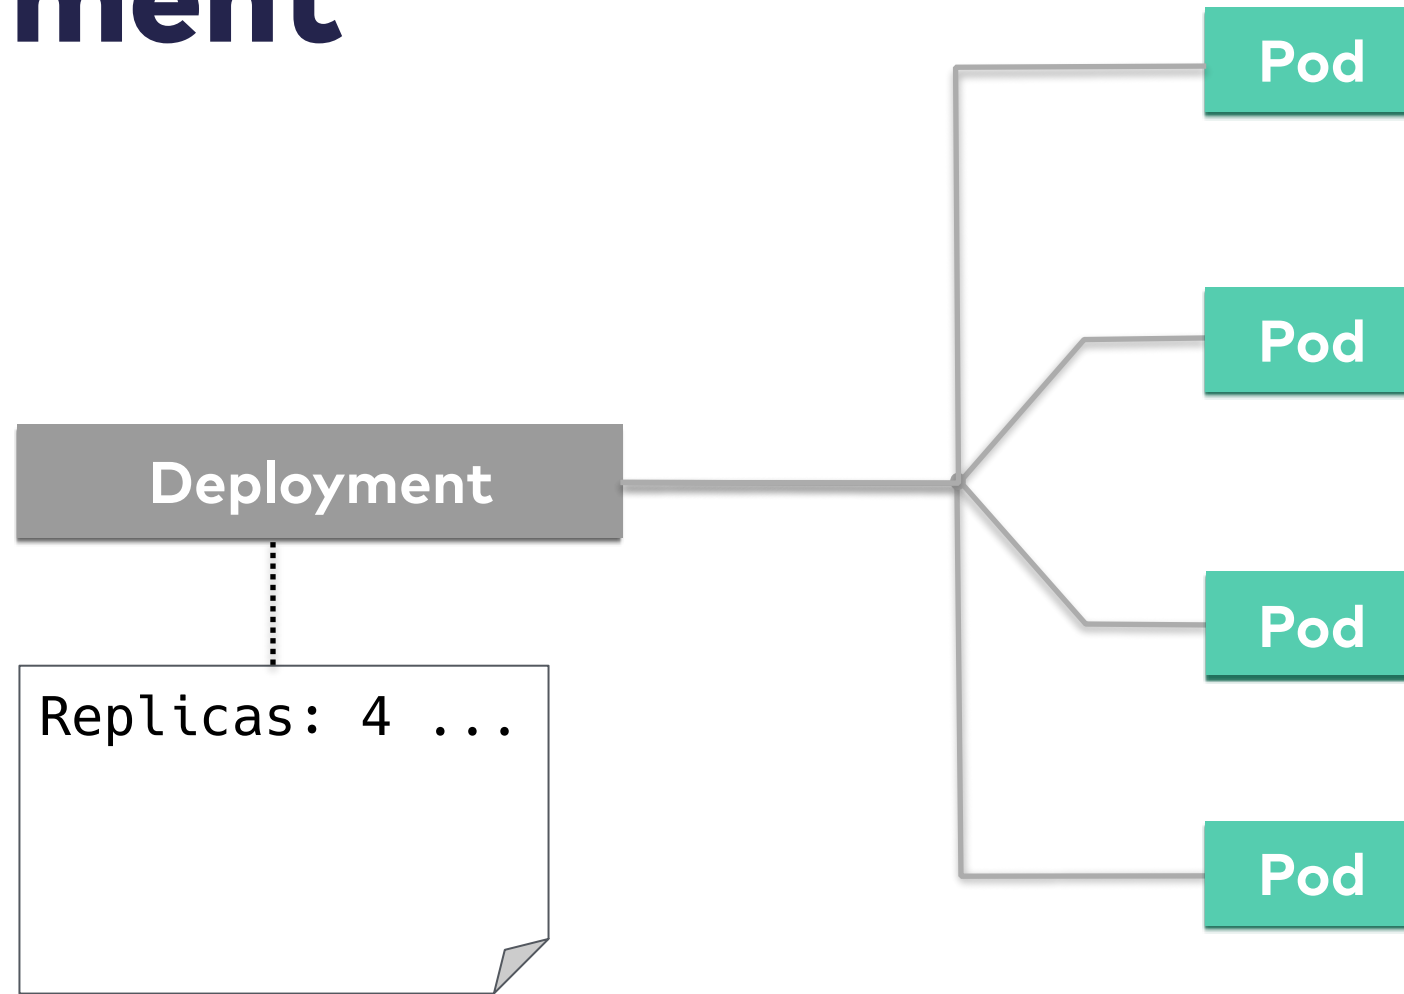
github.com/JoergM/kubernetes_workshop_demos/pods

Deployments

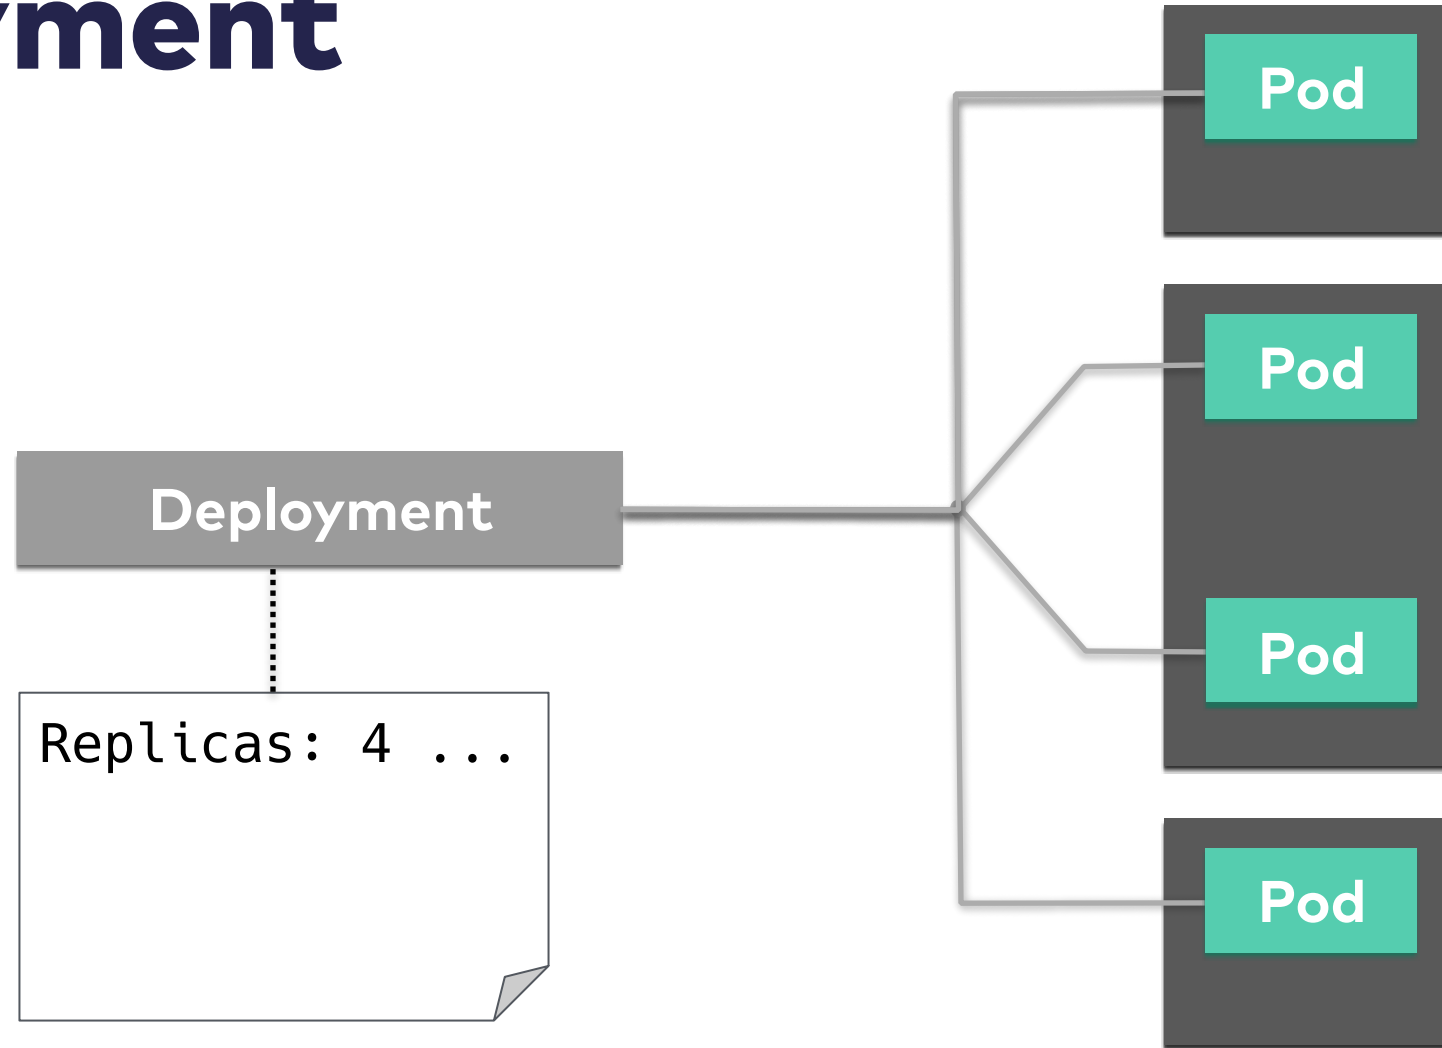
Deployment

- Declares a state of Pods
- Is used for scaling up N instances of the same pod
- Is used to deploy old or new revisions of a pod

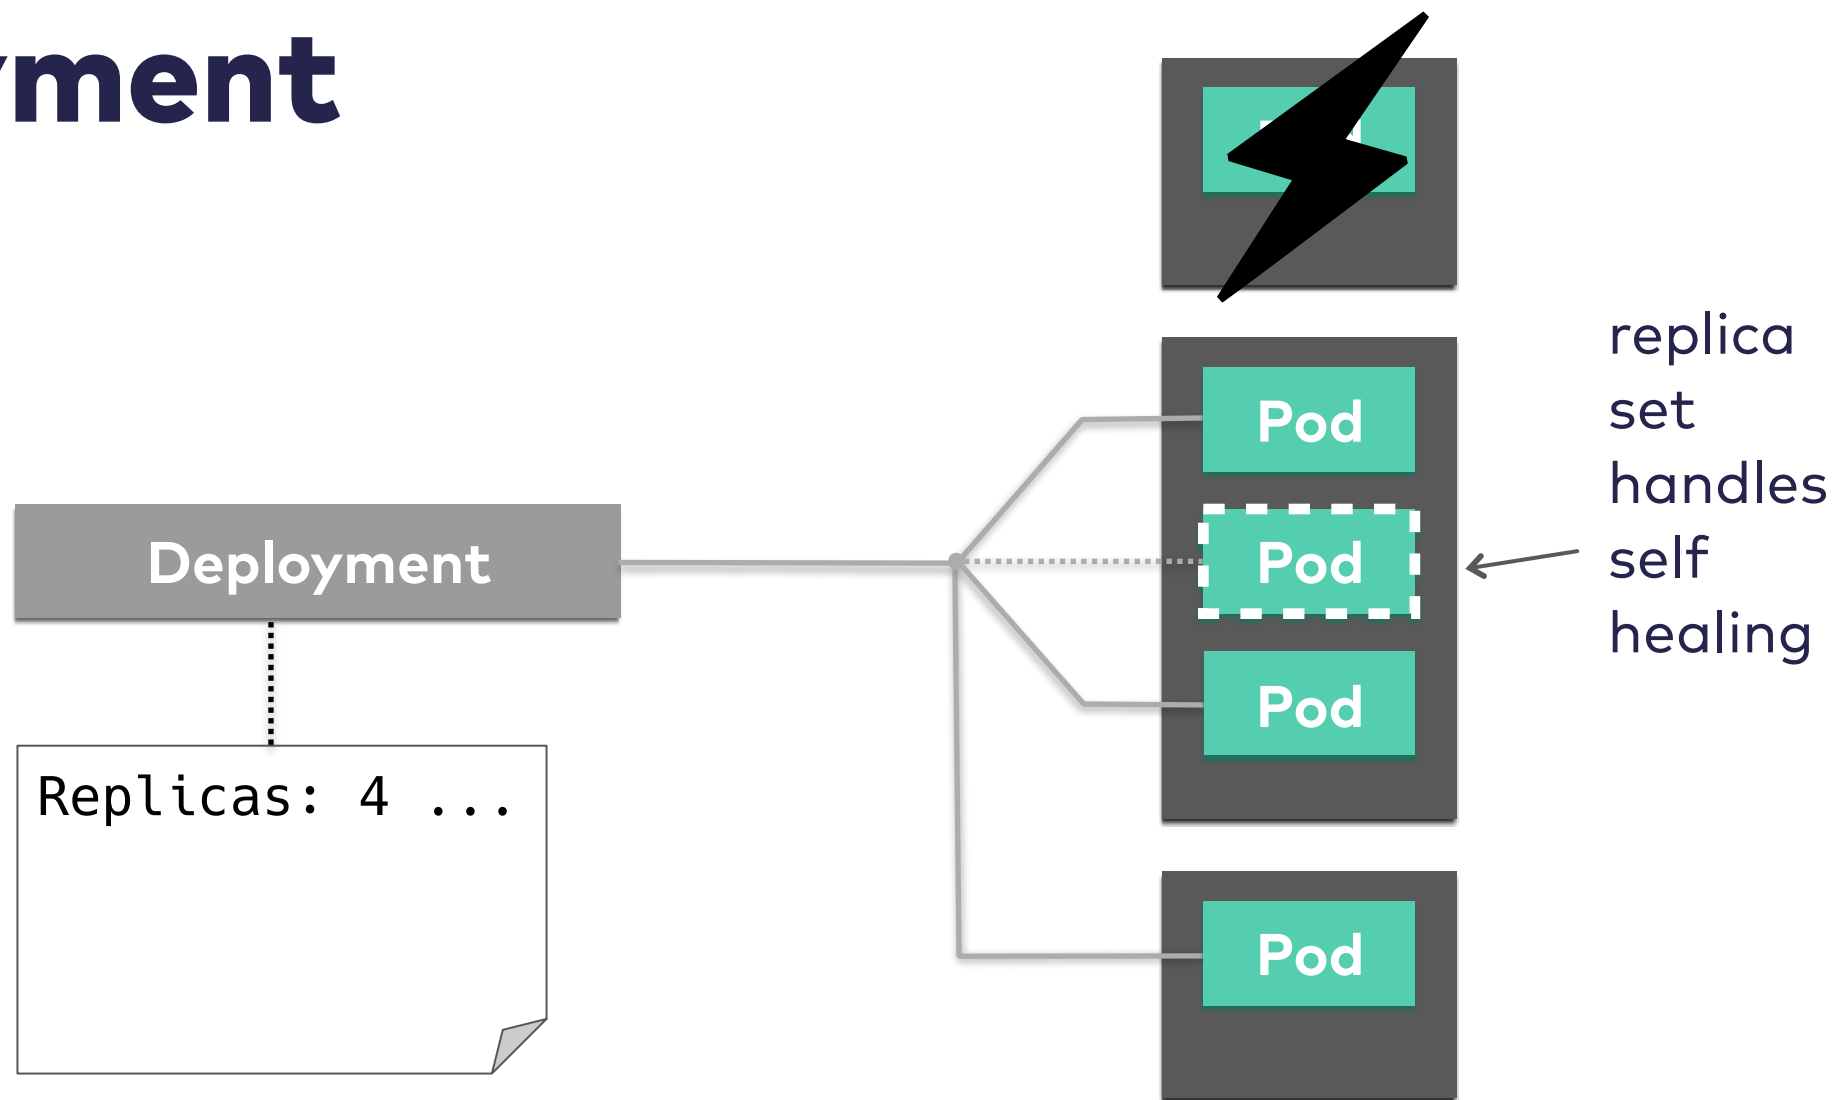
Deployment



Deployment



Deployment




Deploying new versions

- Rolling update
- Recreate
- Blue Green
- Canary


Rolling Update

- Default variant
- No service downtime
- Both versions get traffic at the same time
- consider setting `maxUnavailable/`
`maxSurge`




Recreate

- Activated setting type
- Involves downtime
- no version conflicts




Prelude Service

- IP and DNS for multiple Pods
- Loadbalancing
- Uses Labels to find Pods




Blue / Green

- No version conflicts
- No downtime
- High resource usage
- Involves custom handling by Switching service labels



Canary

- Slowly testing new versions
- No Downtime
- Both versions get traffic at the same time
- Some custom handling of multiple deployments




DEMO

[**github.com/JoergM/kubernetes_workshop_demos/deployments**](https://github.com/JoergM/kubernetes_workshop_demos/deployments)

Services

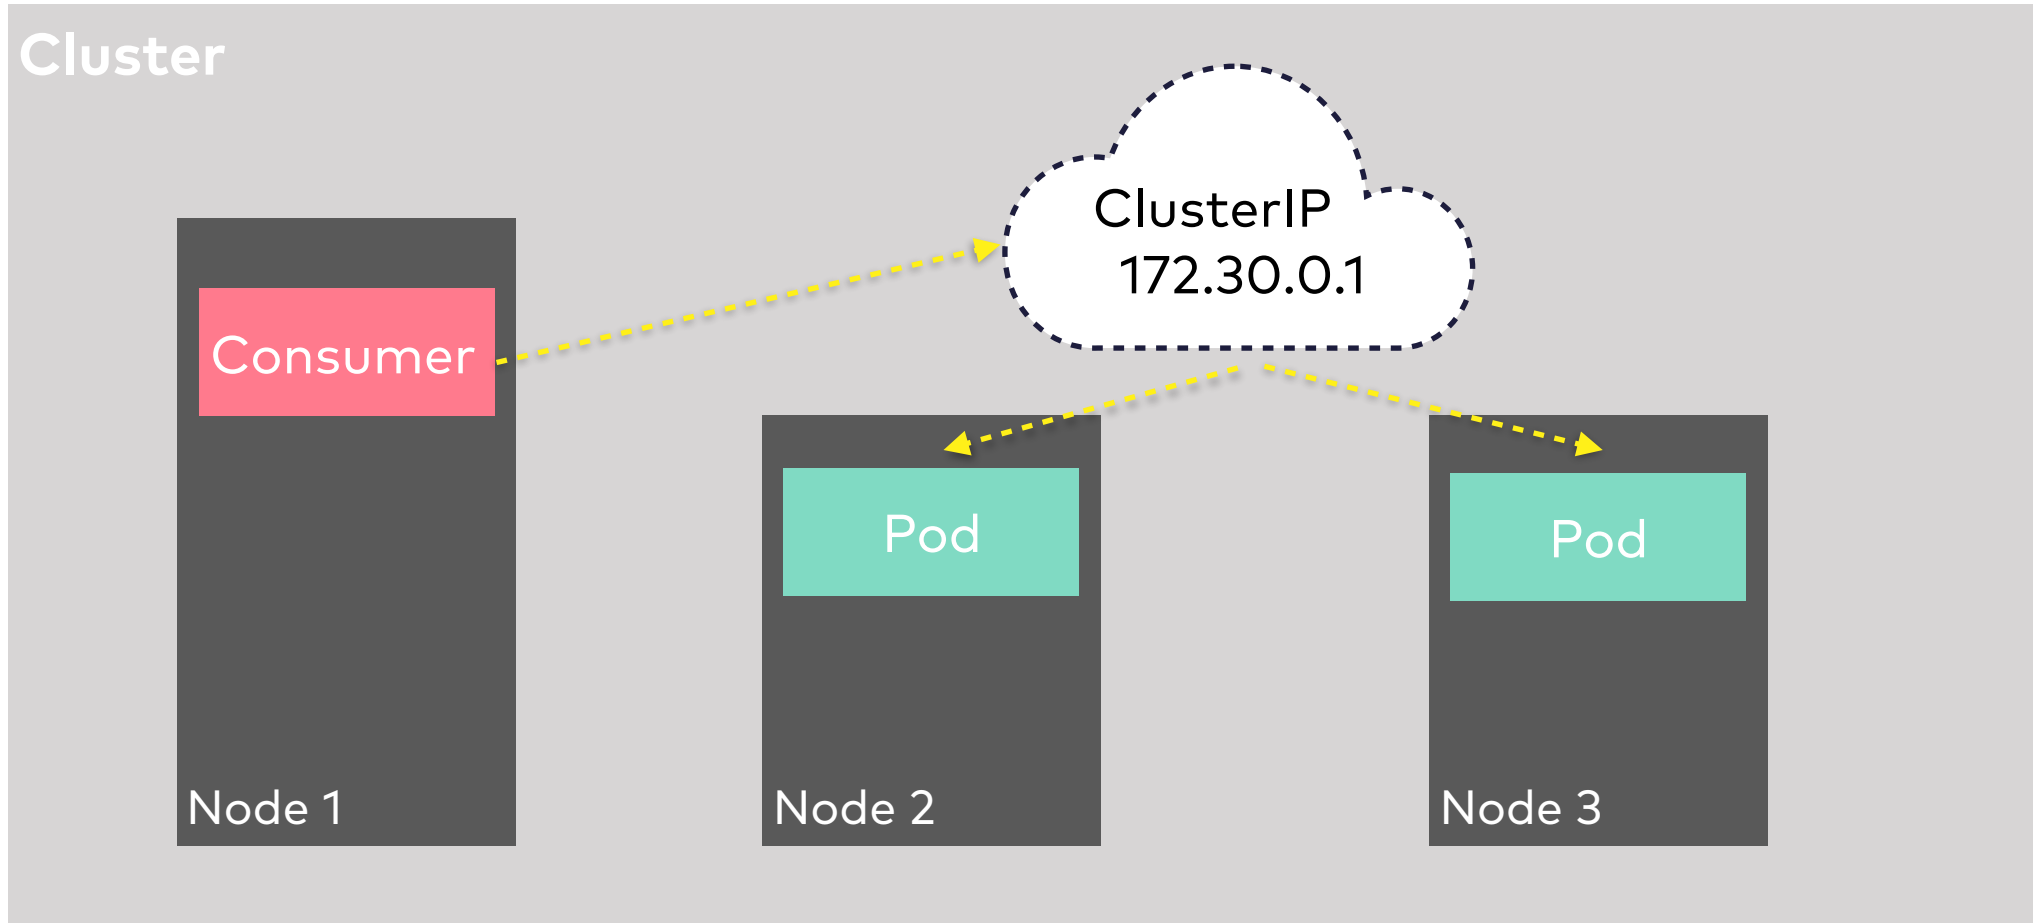
Service Overview



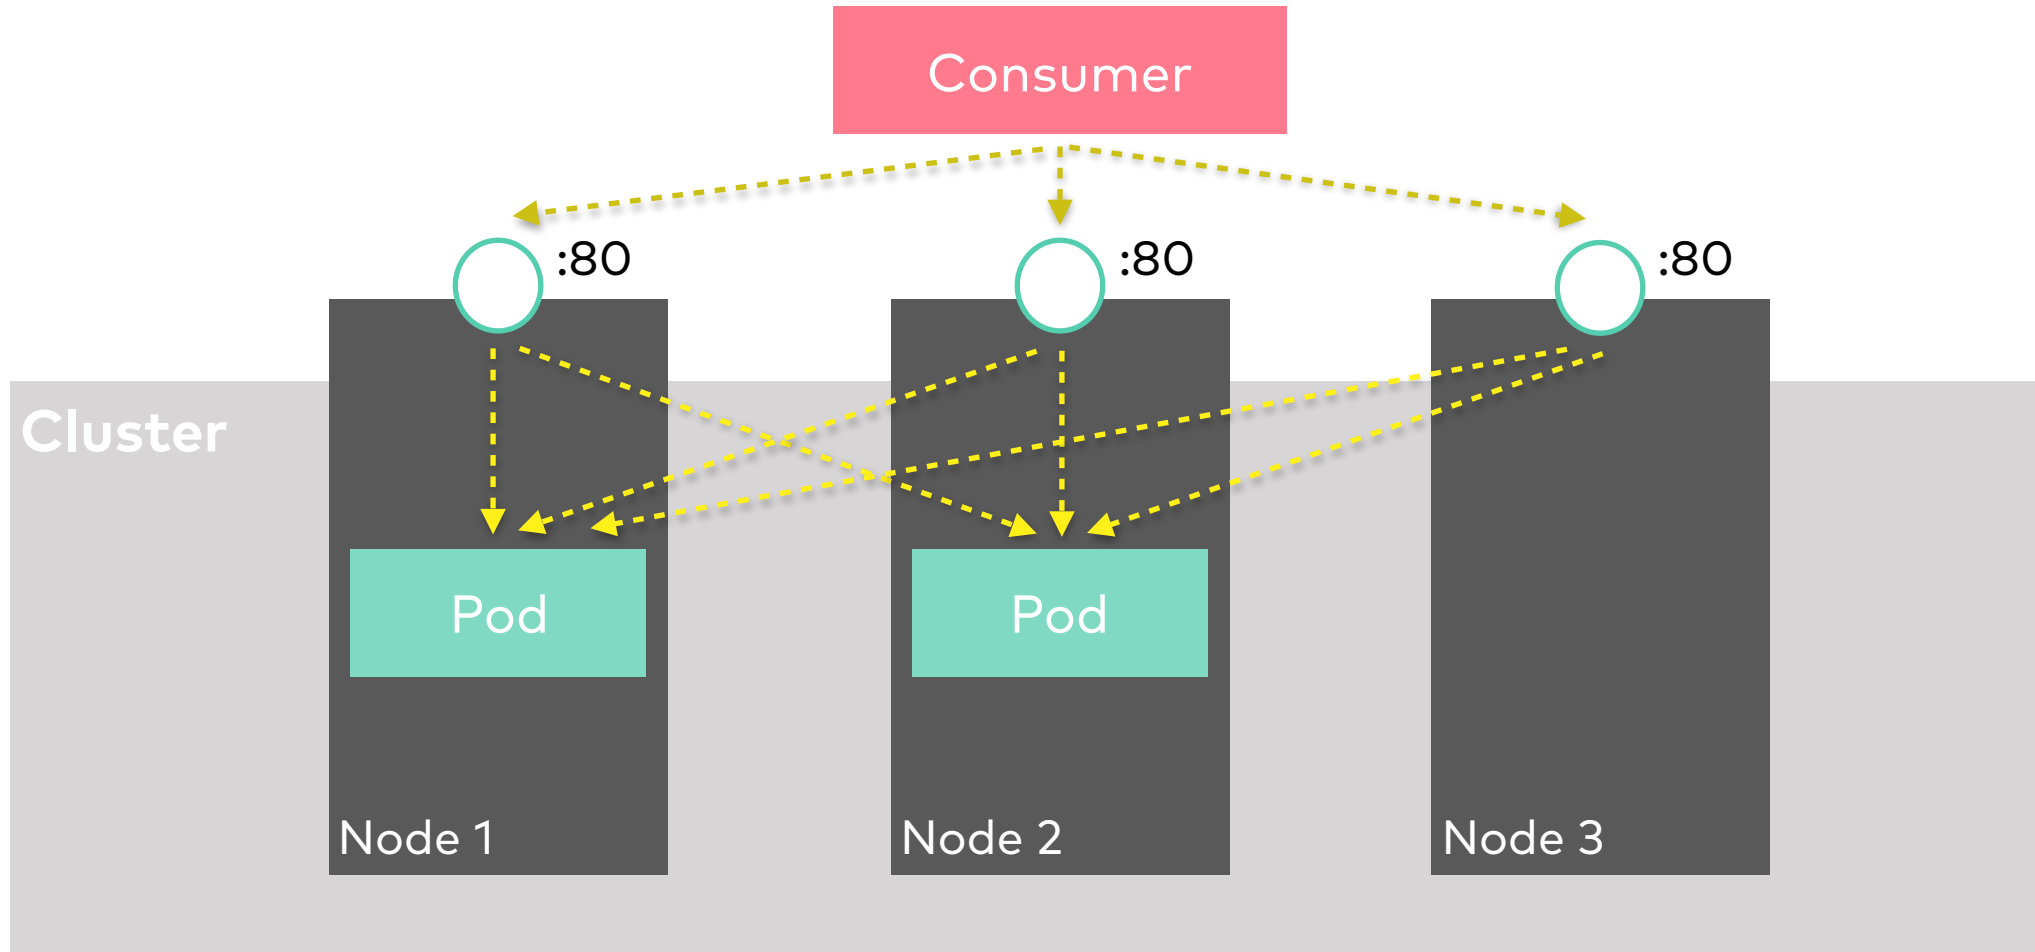
Service

- Is an abstraction which defines a logical set of pods and a policy by which to access them
- Usually represents a micro-service
- Different types of services possible
- Discovery inside Cluster via DNS
- It's not a physical LoadBalancer (more later)

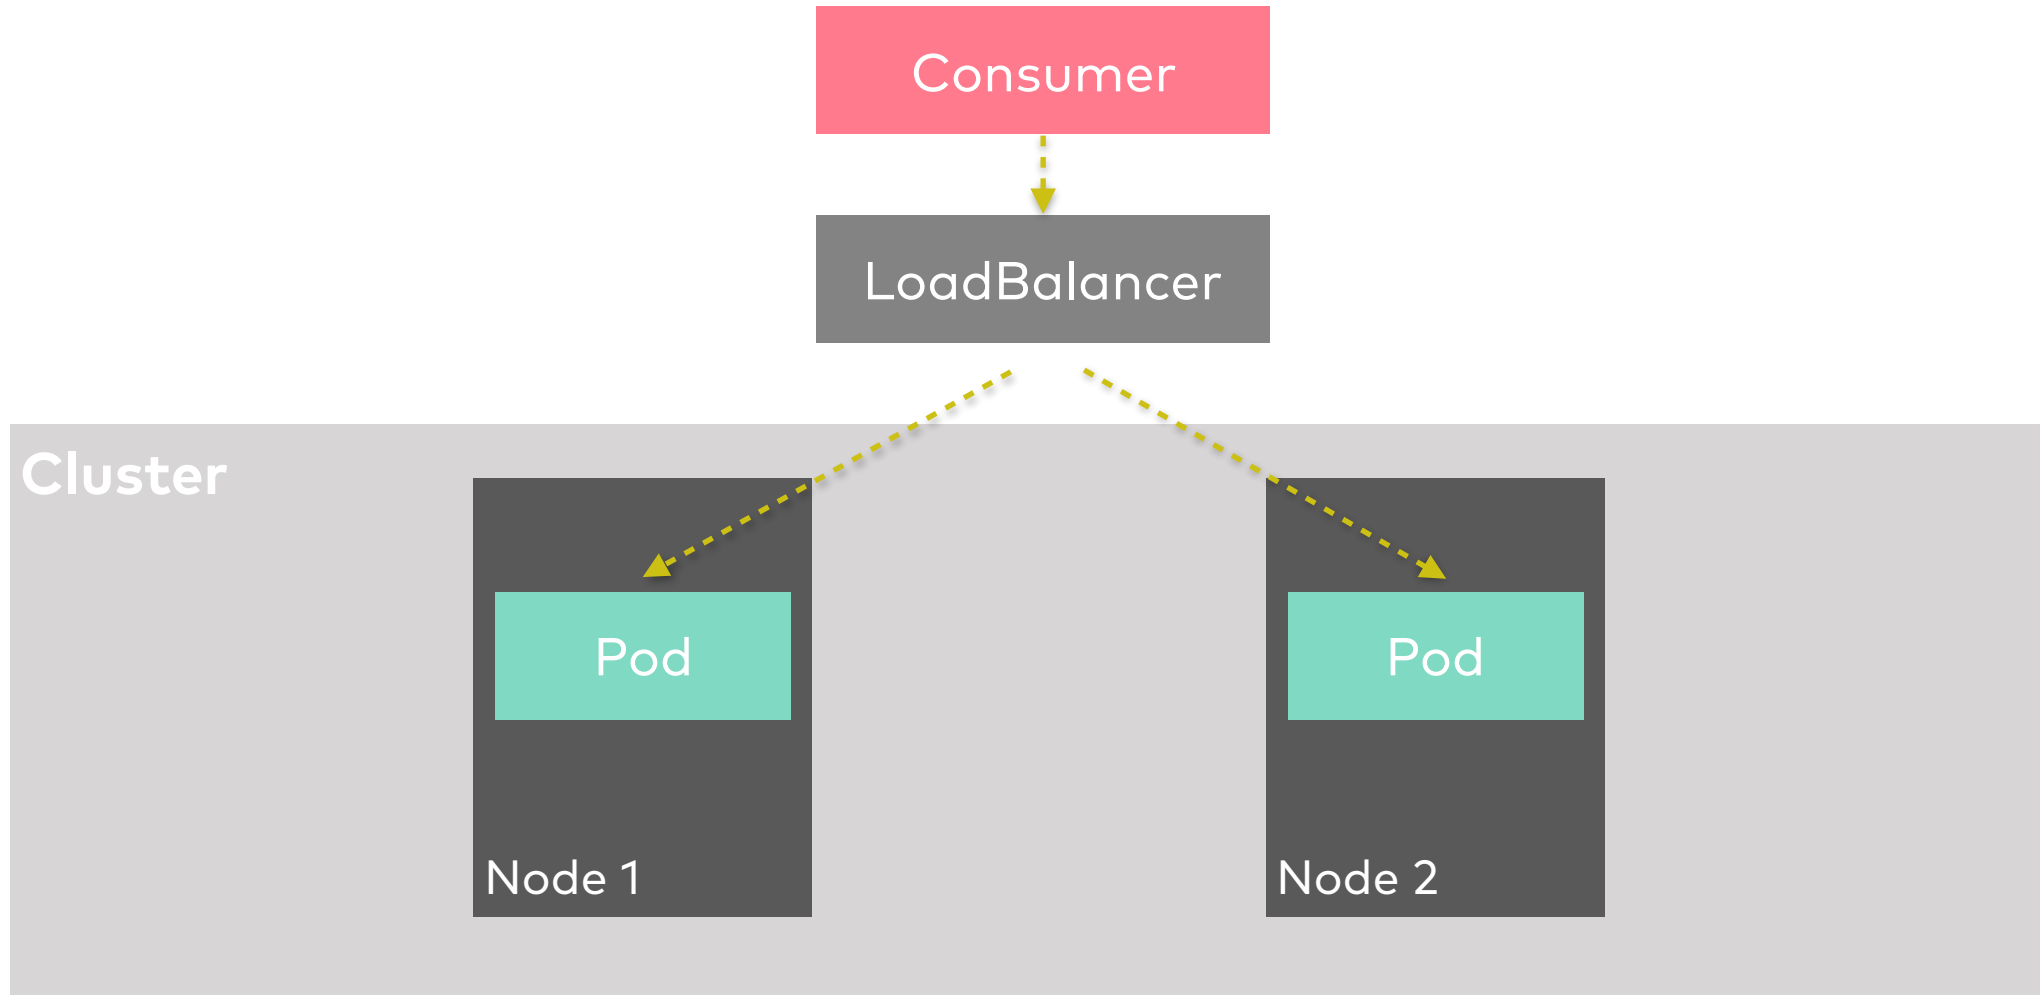
Service type ClusterIP




Service type NodePort



Service type LoadBalancer



Service type External...



DEMO

[**github.com/JoergM/kubernetes_workshop_demos/services**](https://github.com/JoergM/kubernetes_workshop_demos/services)

Configuration

Config Maps

- Provide Pods with configuration data
- from
 - literal values
 - files
 - directories

Config Maps


- In Pods as
 - environment variables
 - files
 - directories

DEMO

github.com/JoergM/kubernetes_workshop_demos/configuration

Ingress

Ingress



Ingress Controller

- **Creates a LoadBalancer service that points to a pod, which runs a reverse proxy (nginx, haproxy, Apache, traefik)**
- **Uses IngressRules to describe which DNS and/or path should point to which service**
- **Always needs a default service**

Ingress controller implementations

- **nginx**
- **traefik**
- **voyager**
- **GCE ingress**
- **Kong**
- **...**


DEMO

[**github.com/JoergM/kubernetes_workshop_demos/ingress**](https://github.com/JoergM/kubernetes_workshop_demos/ingress)

Jobs


Jobs

- Running pods until **completion**
- Like deployment for long-running pods
- supports parallelism



CronJob

- Regularly starting jobs
- Follows typical Cron patterns:
 - 0 12 * * 1-5
 - (weekdays at noon)




DEMO

github.com/JoergM/kubernetes_workshop_demos/jobs

Persistence

Persistence Overview



Persistent Volume Claims

- User requesting Storage
- Used in pod as volume
- Survives pod recreation
- Certain Size (e.g. 5Gi)
- Certain class (e.g. SSD)
- Will be matched to persistent volumes

Persistent Volumes

- Defines a real volume of a certain size (e.g. 5Gi)
- Can be created upfront
- Lots of implementations:
 - GCEPersistenceDisk
 - HostPath
 - AWS EBS
 - NFS

Dynamic provisioning

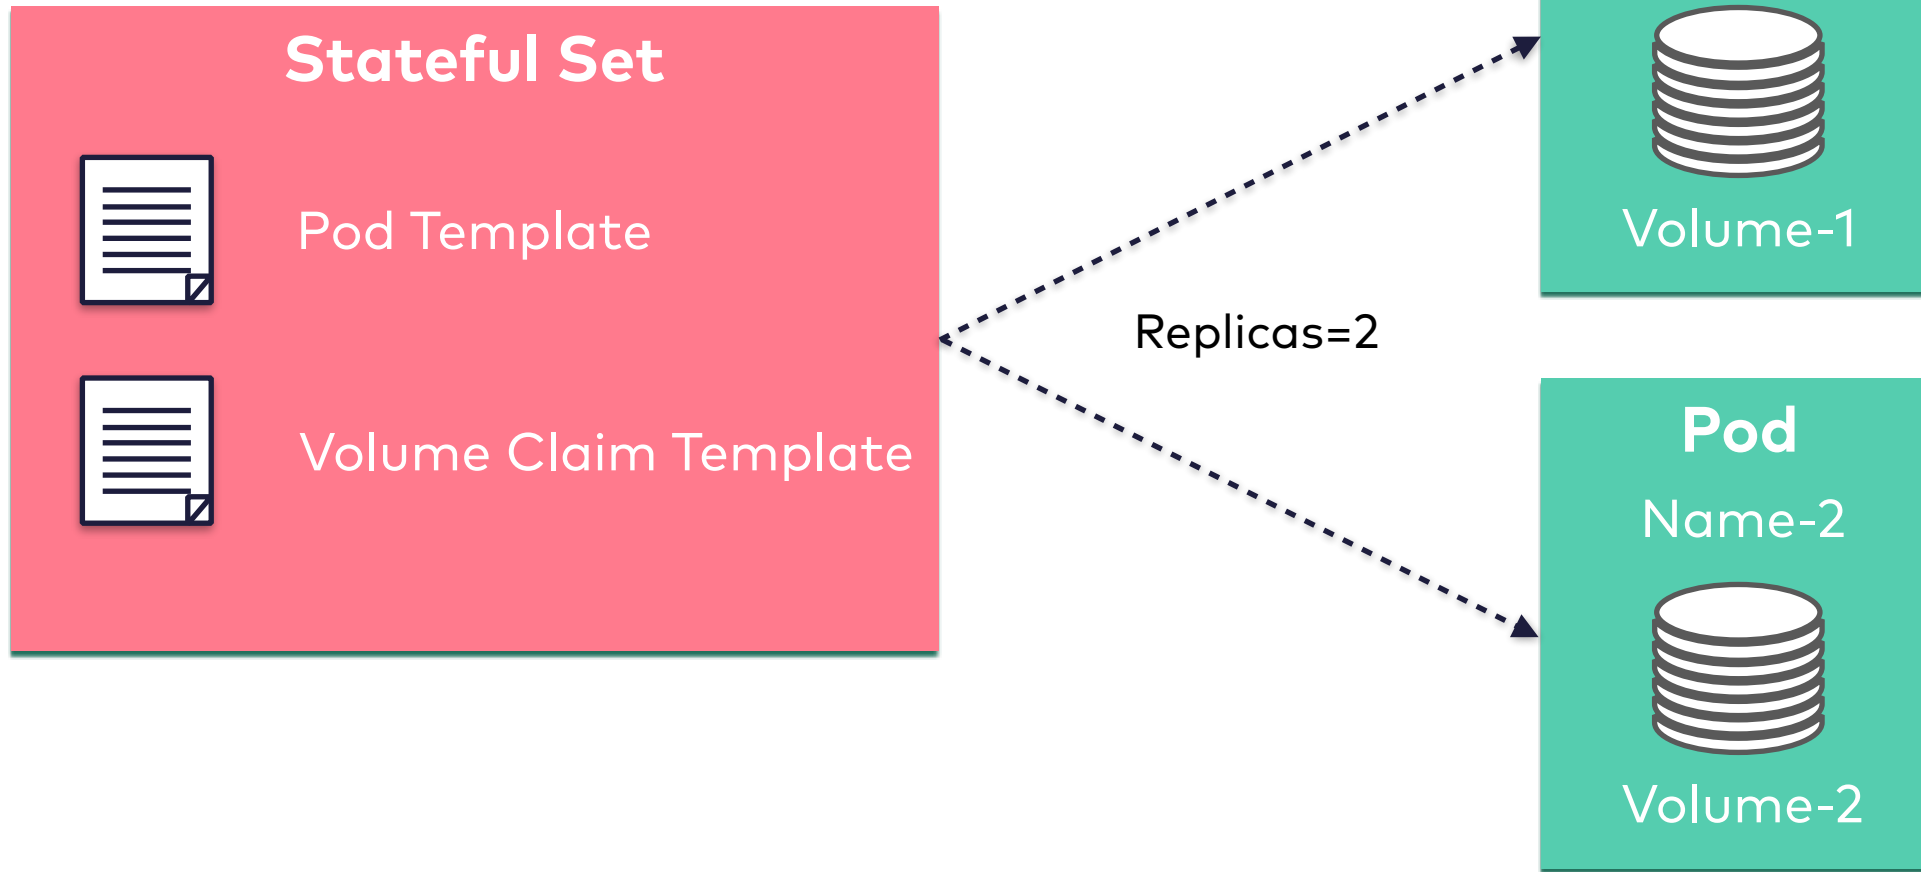
- Creating Volumes based on Claims
- Requires dynamic way of creating volumes and mounting to nodes
- e.g. AWS EBS, GCE Persistent Disk
- Custom provisioners for can be created too

DEMO

[**github.com/JoergM/kubernetes_workshop_demos/
persistent_volumes**](https://github.com/JoergM/kubernetes_workshop_demos/persistent_volumes)

Stateful Sets

Overview



Stateful Sets

- Like Deployment but with other guarantees
- Each Replica has always the same name (and DNS)
- Replica and Volume Claim always come together
- Most parameters not changeable after creation

Stateful Sets - usages

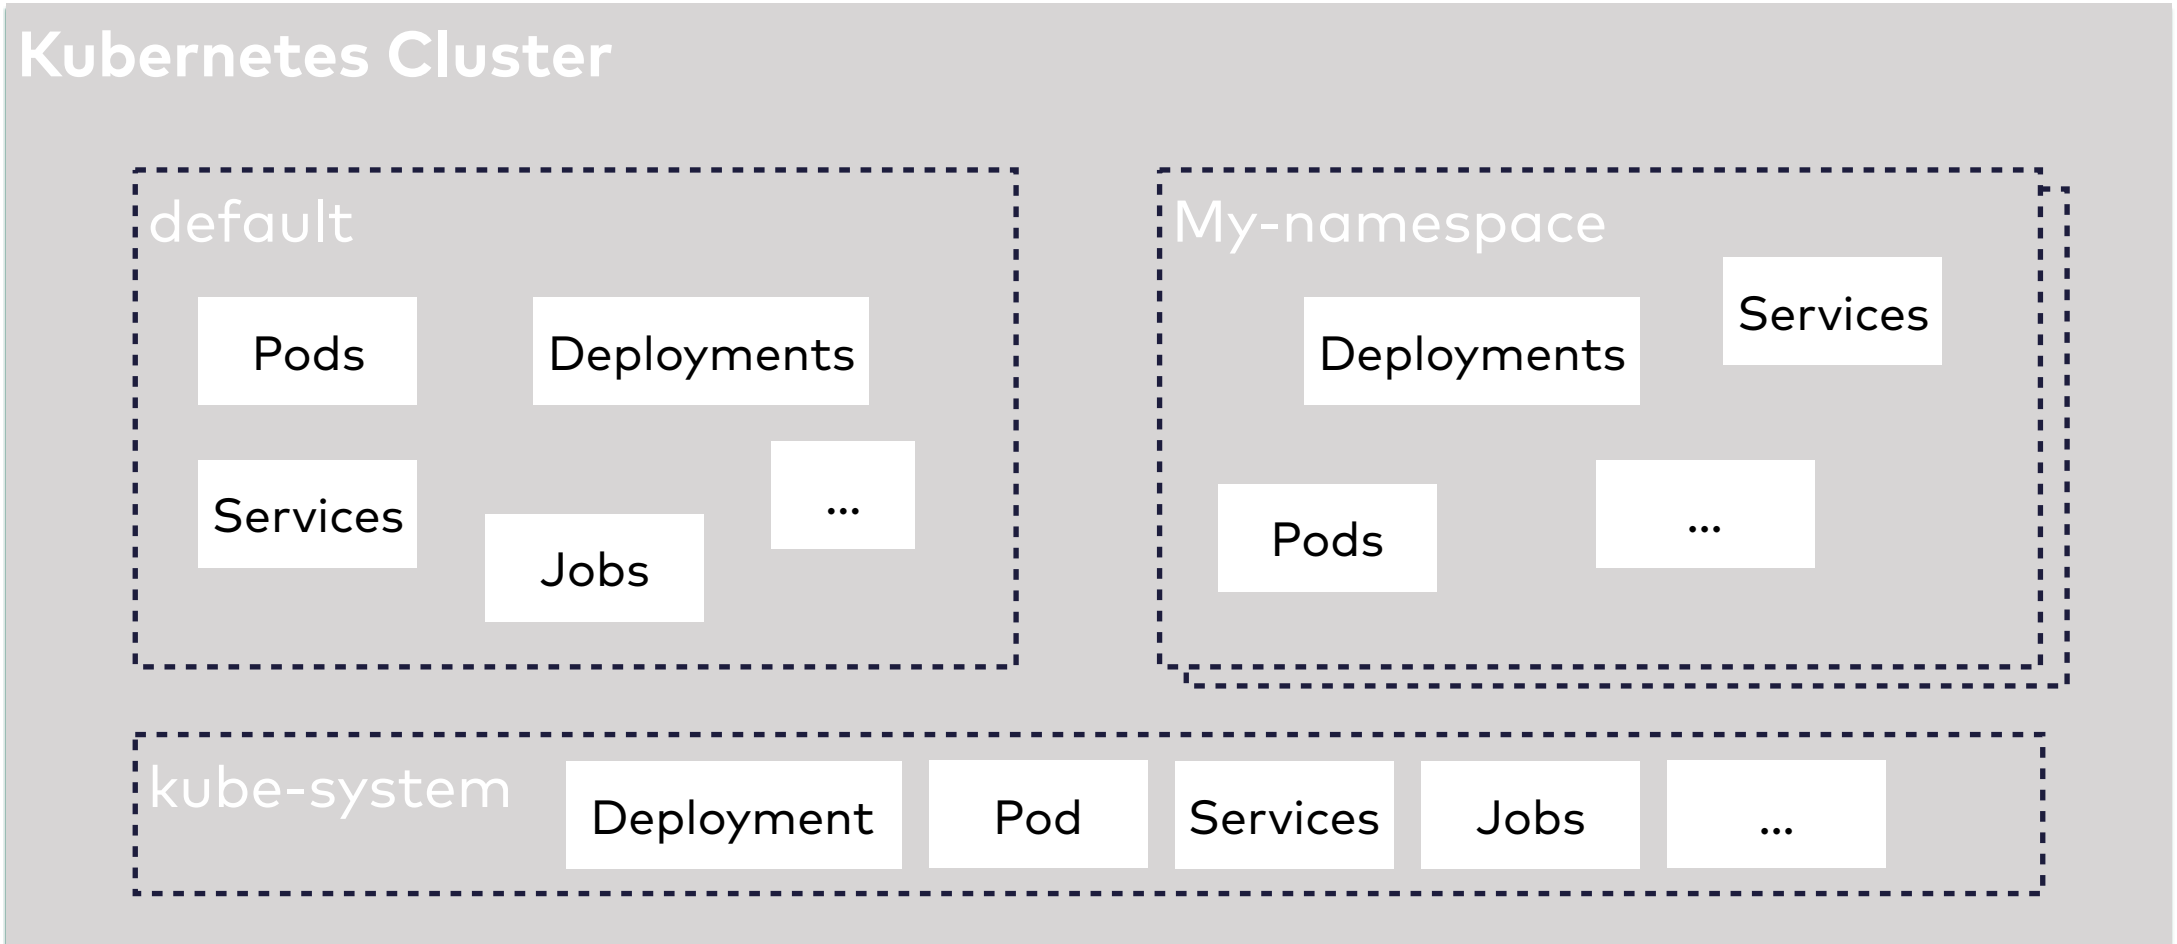
- All kind of software that builds a cluster, but needs certain guarantees
- e.g. Databases with fixed Follower-Leader specification
 - MongoDB
 - Zookeeper
 - Postgresql
- Do not use if not necessary!

DEMO

[**github.com/JoergM/kubernetes_workshop_demos/stateful_sets**](https://github.com/JoergM/kubernetes_workshop_demos/stateful_sets)

Namespaces

Namespaces overview



Namespaces

- Scope for names of objects
- Hook for service accounts and network policies
- Isolation level depends on your installation
- Not all objects are in namespaces (esp. low level like nodes or persistent volumes)

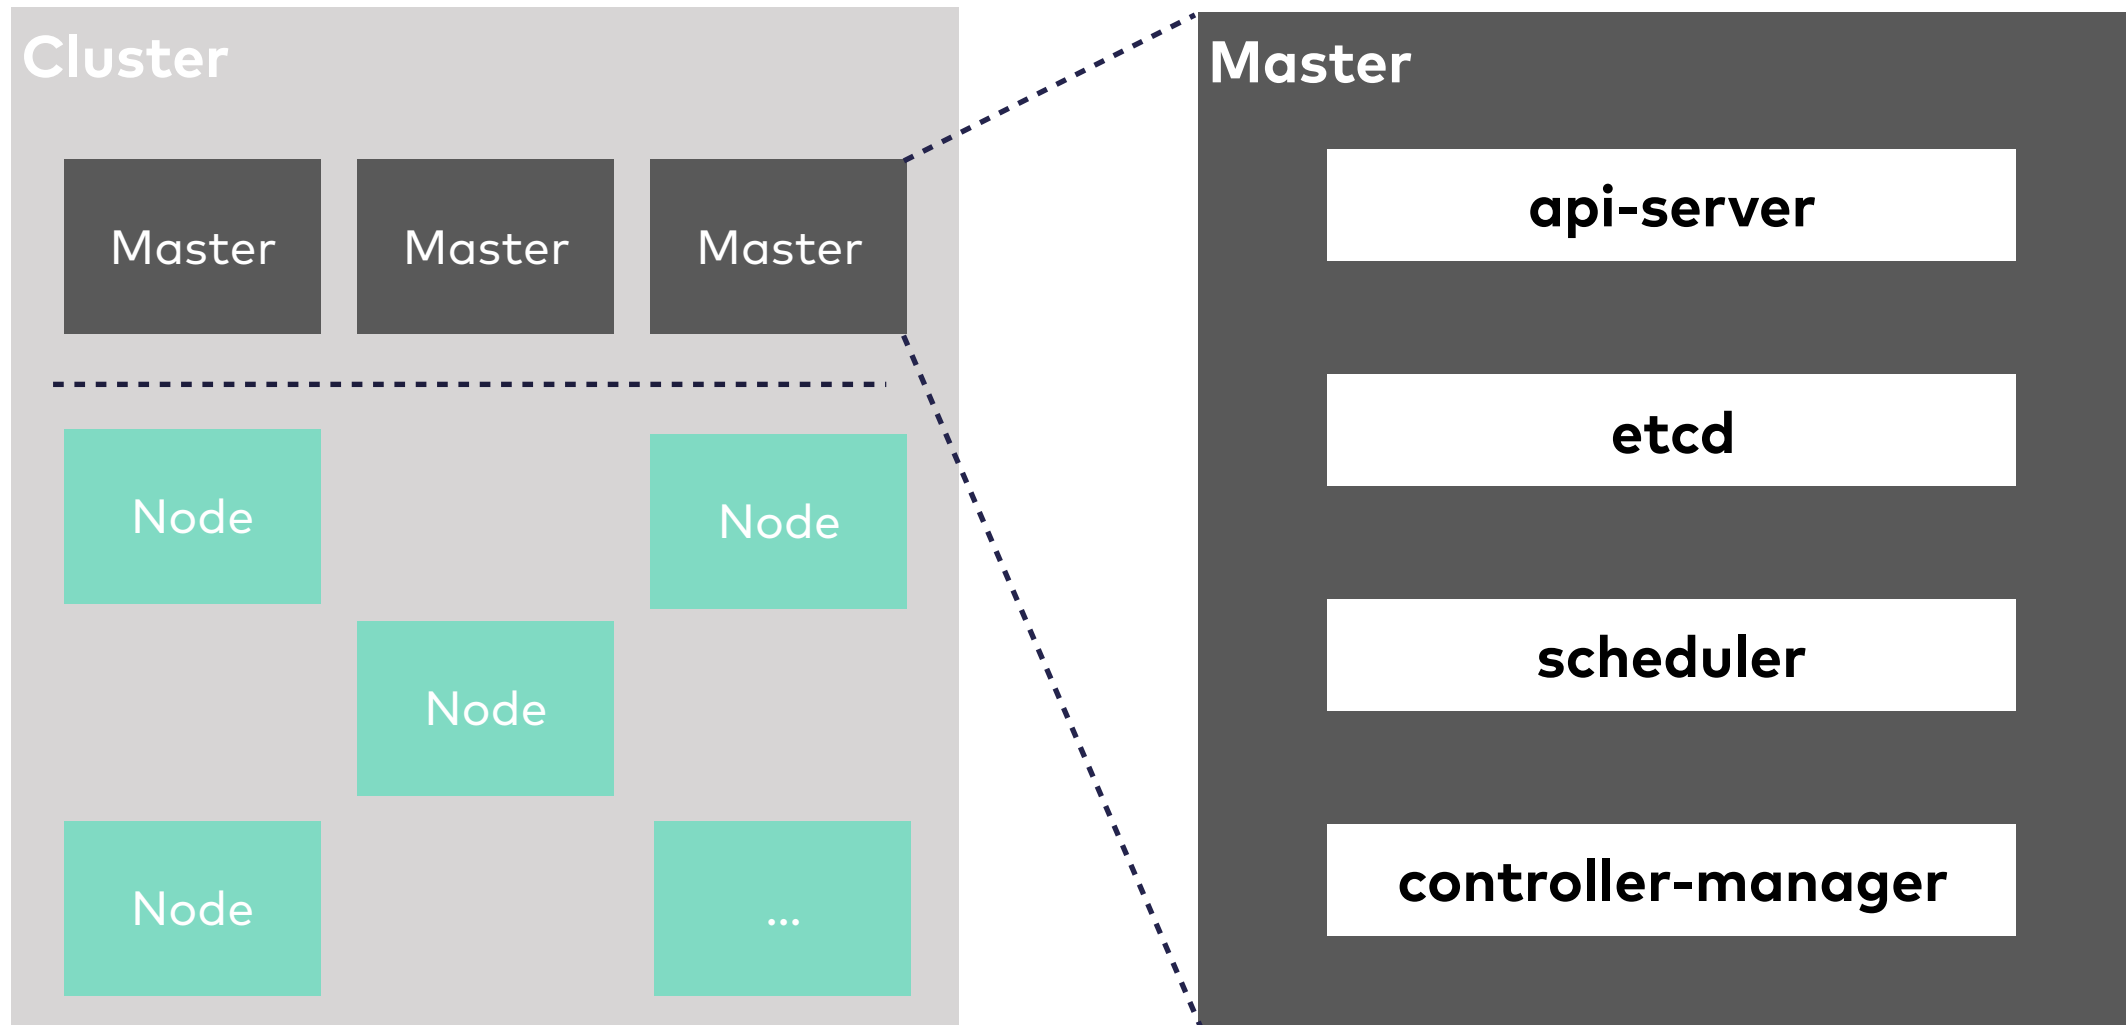
DEMO

github.com/JoergM/kubernetes_workshop_demos/namespaces



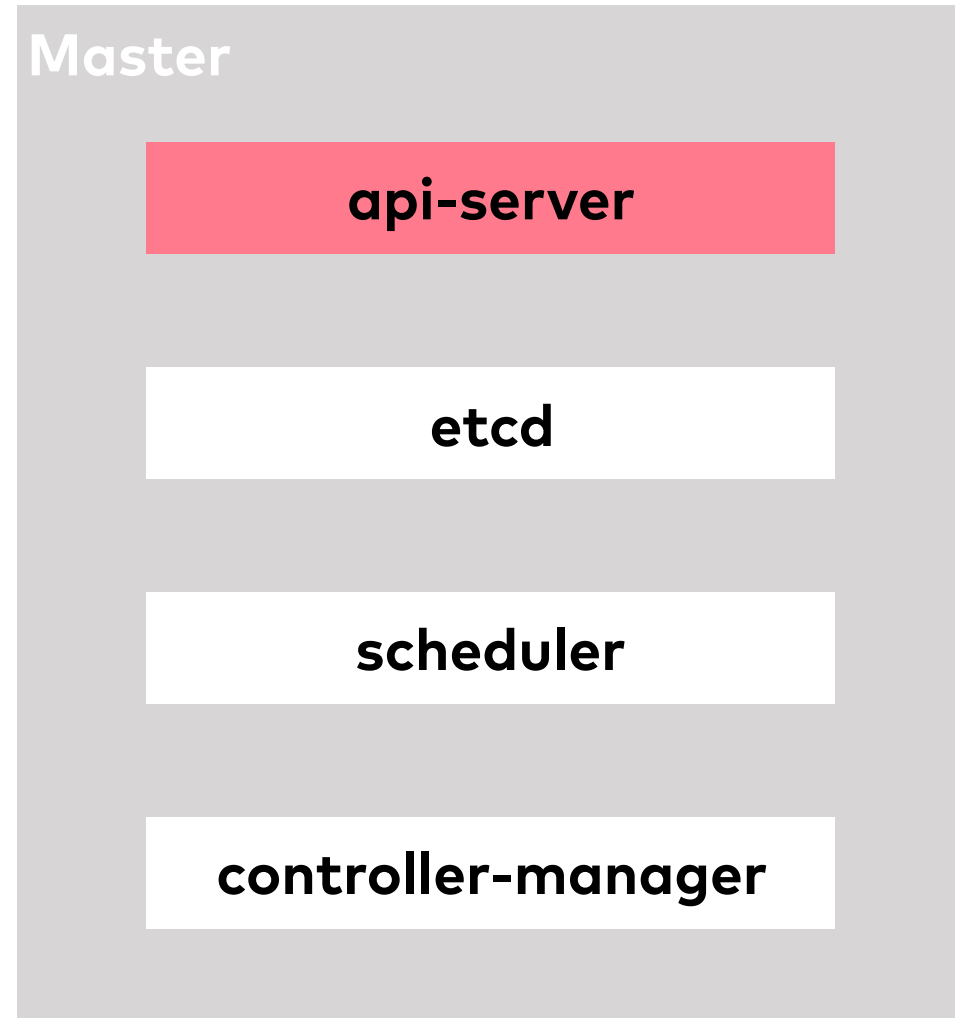
Internal architecture

Master Components



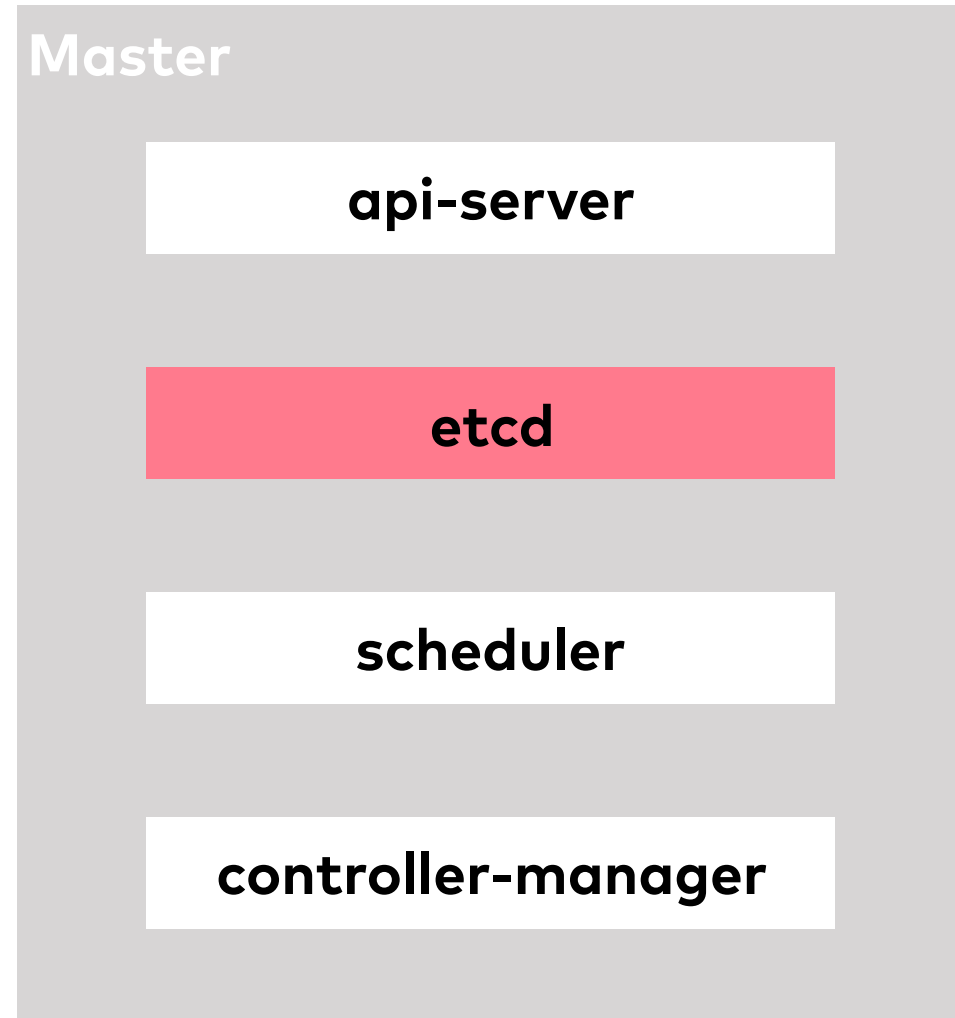
API Server

- Entry point for all interactions
- Stores desired state into etcd
- available from outside and inside cluster



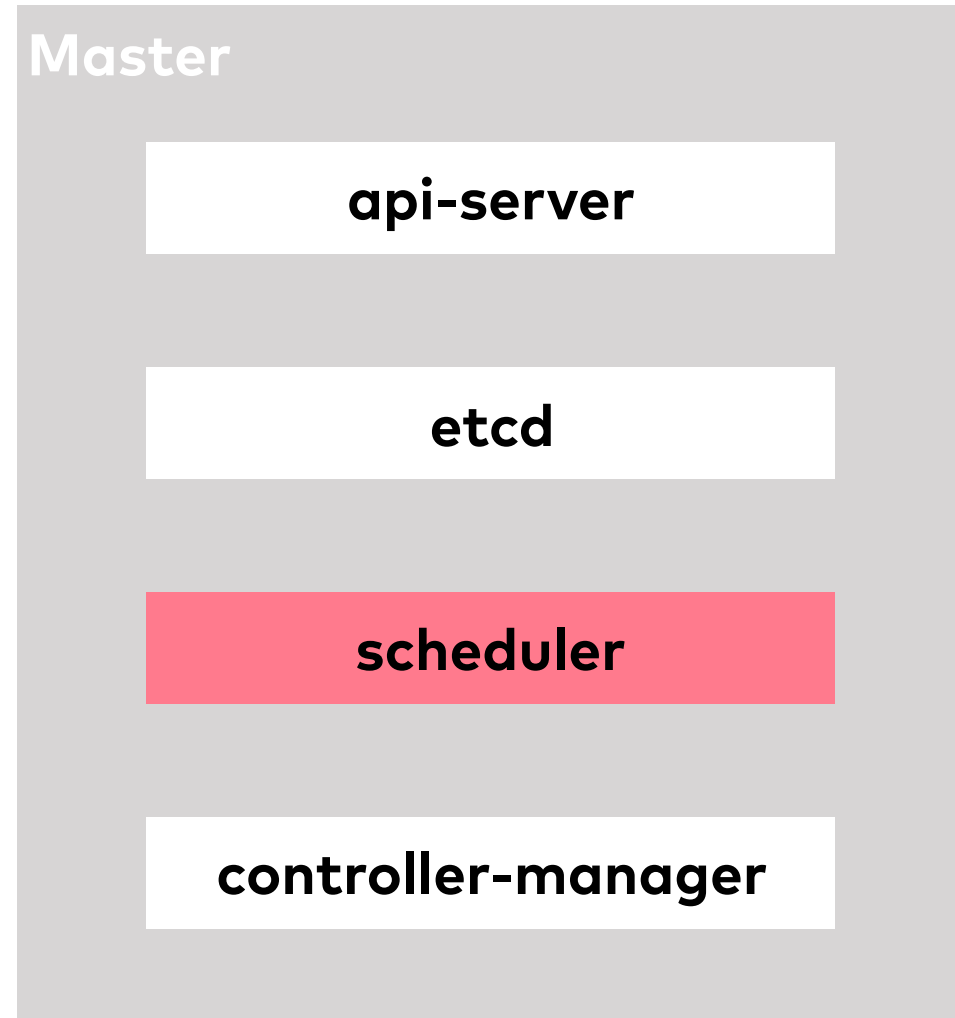
etcd

- consensus based distributed key value database
- interaction only via api-server



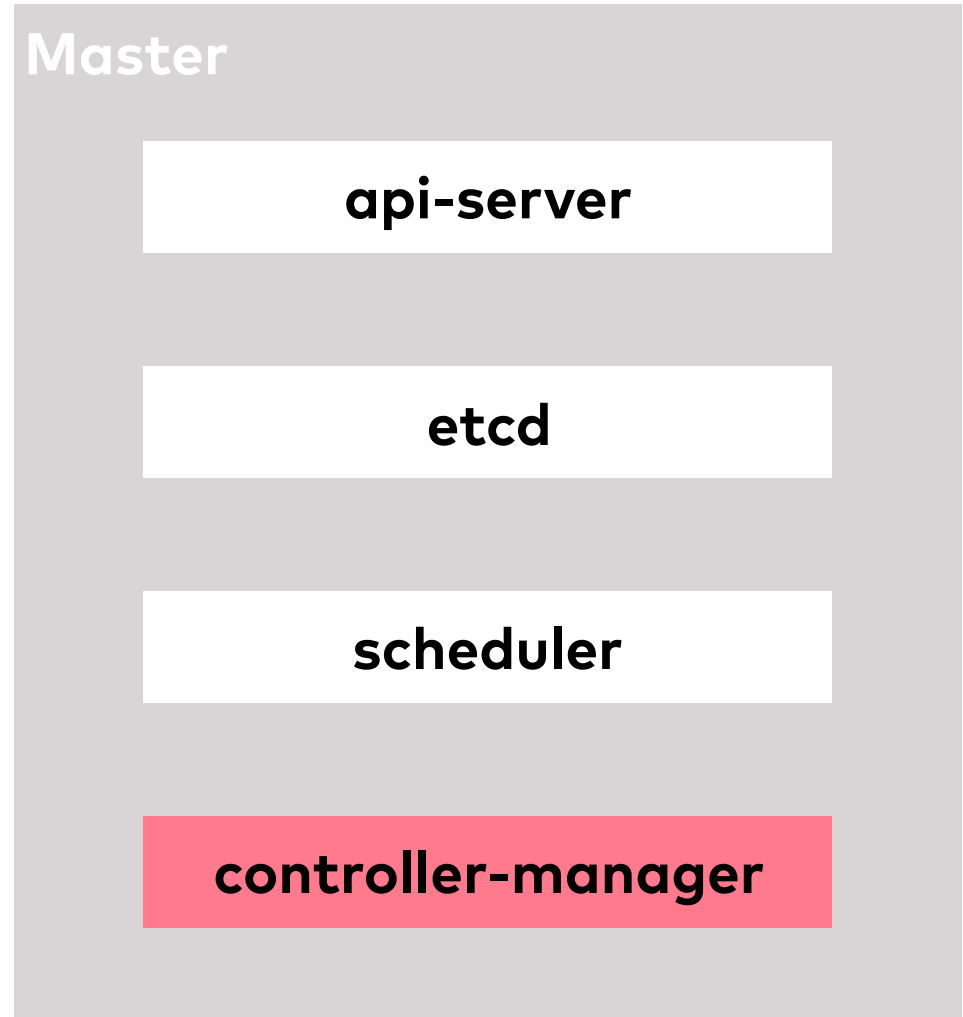
Scheduler

- Watches newly created pods and assigns them to nodes
- Lots of criterias
 - resource requirements
 - load
 - specific constraints



Controller-Manager

- **Manages / runs the controllers responsible for certain tasks in Kubernetes**




Kubernetes API – Controller

- Watch the Api Server for changes
- perform Operations on changes
- Creation/ Deletion or Update on other API objects
- Running a reconciliation loop

Controller examples

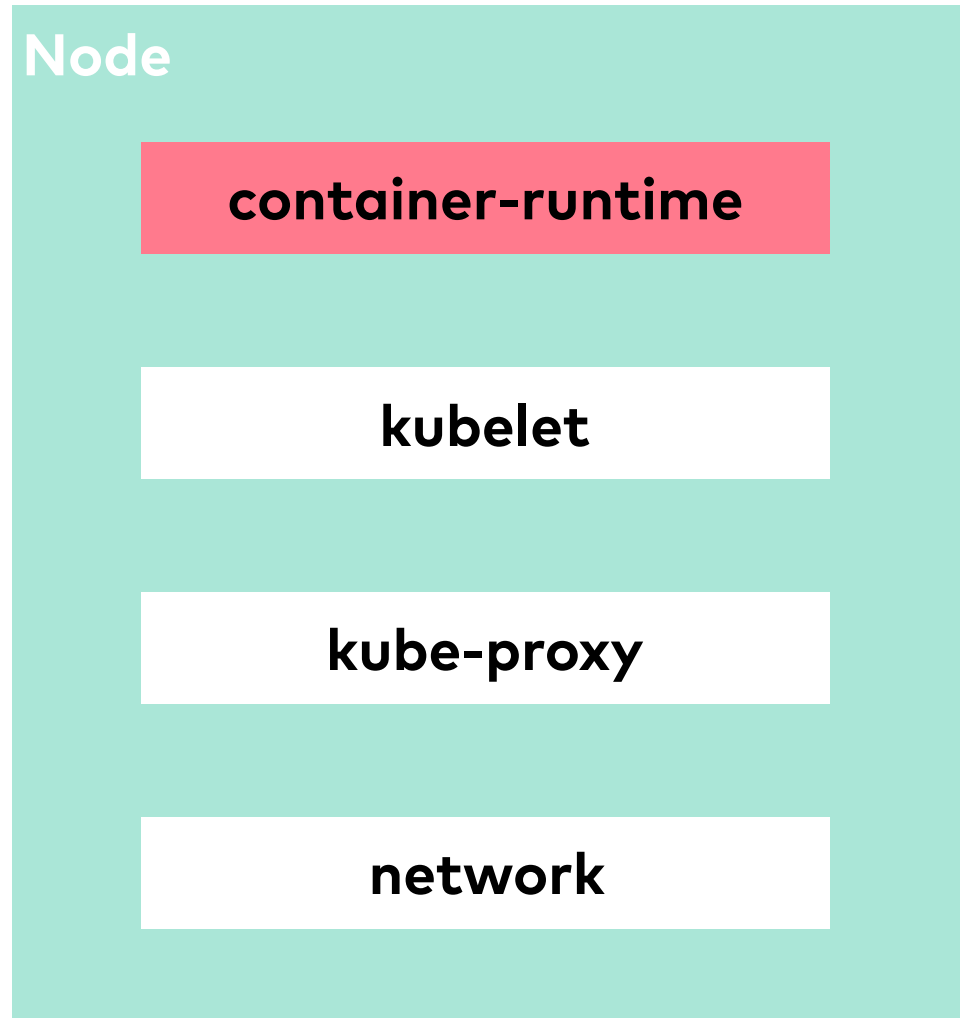
- Node Controller
- Replication Controller
- Cloud Volume Controller
- DNS Controller
- (your controller)

Node Components



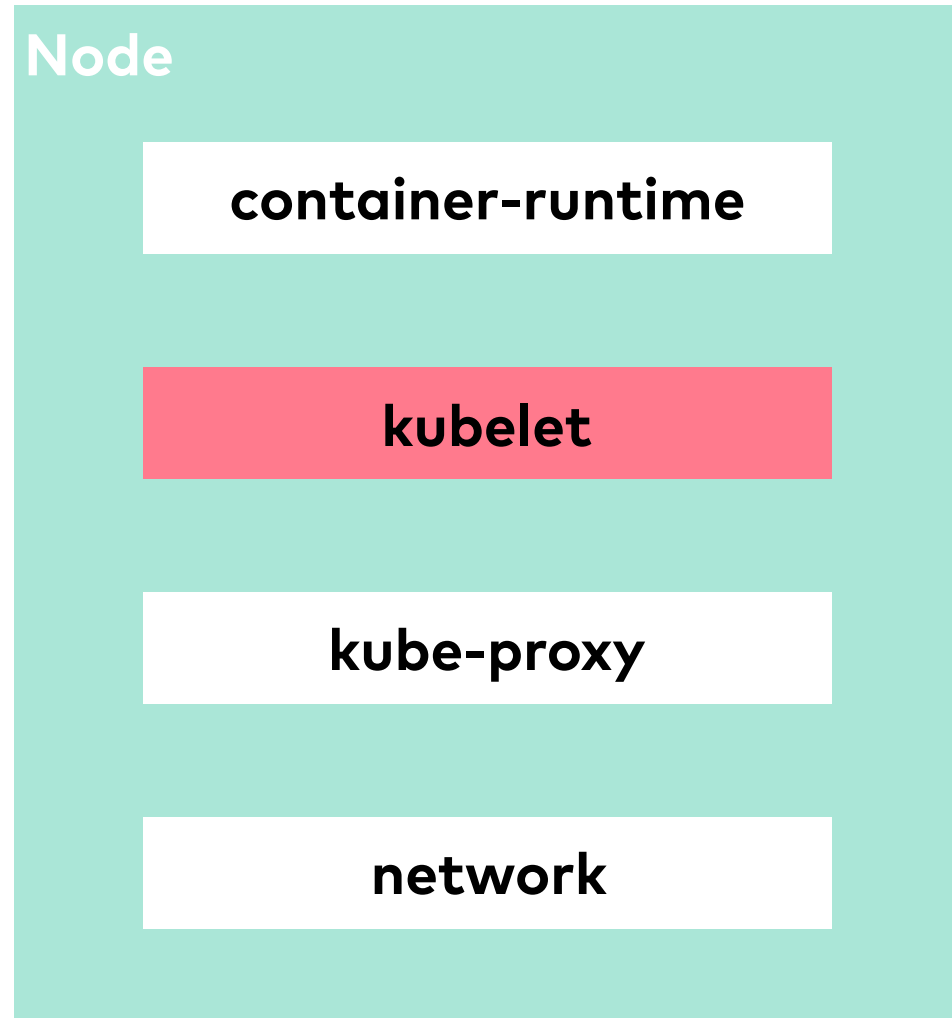
Container runtime

- Component to run containers on nodes
- Usually Docker
- Can be other implementation (e.g. rkt)

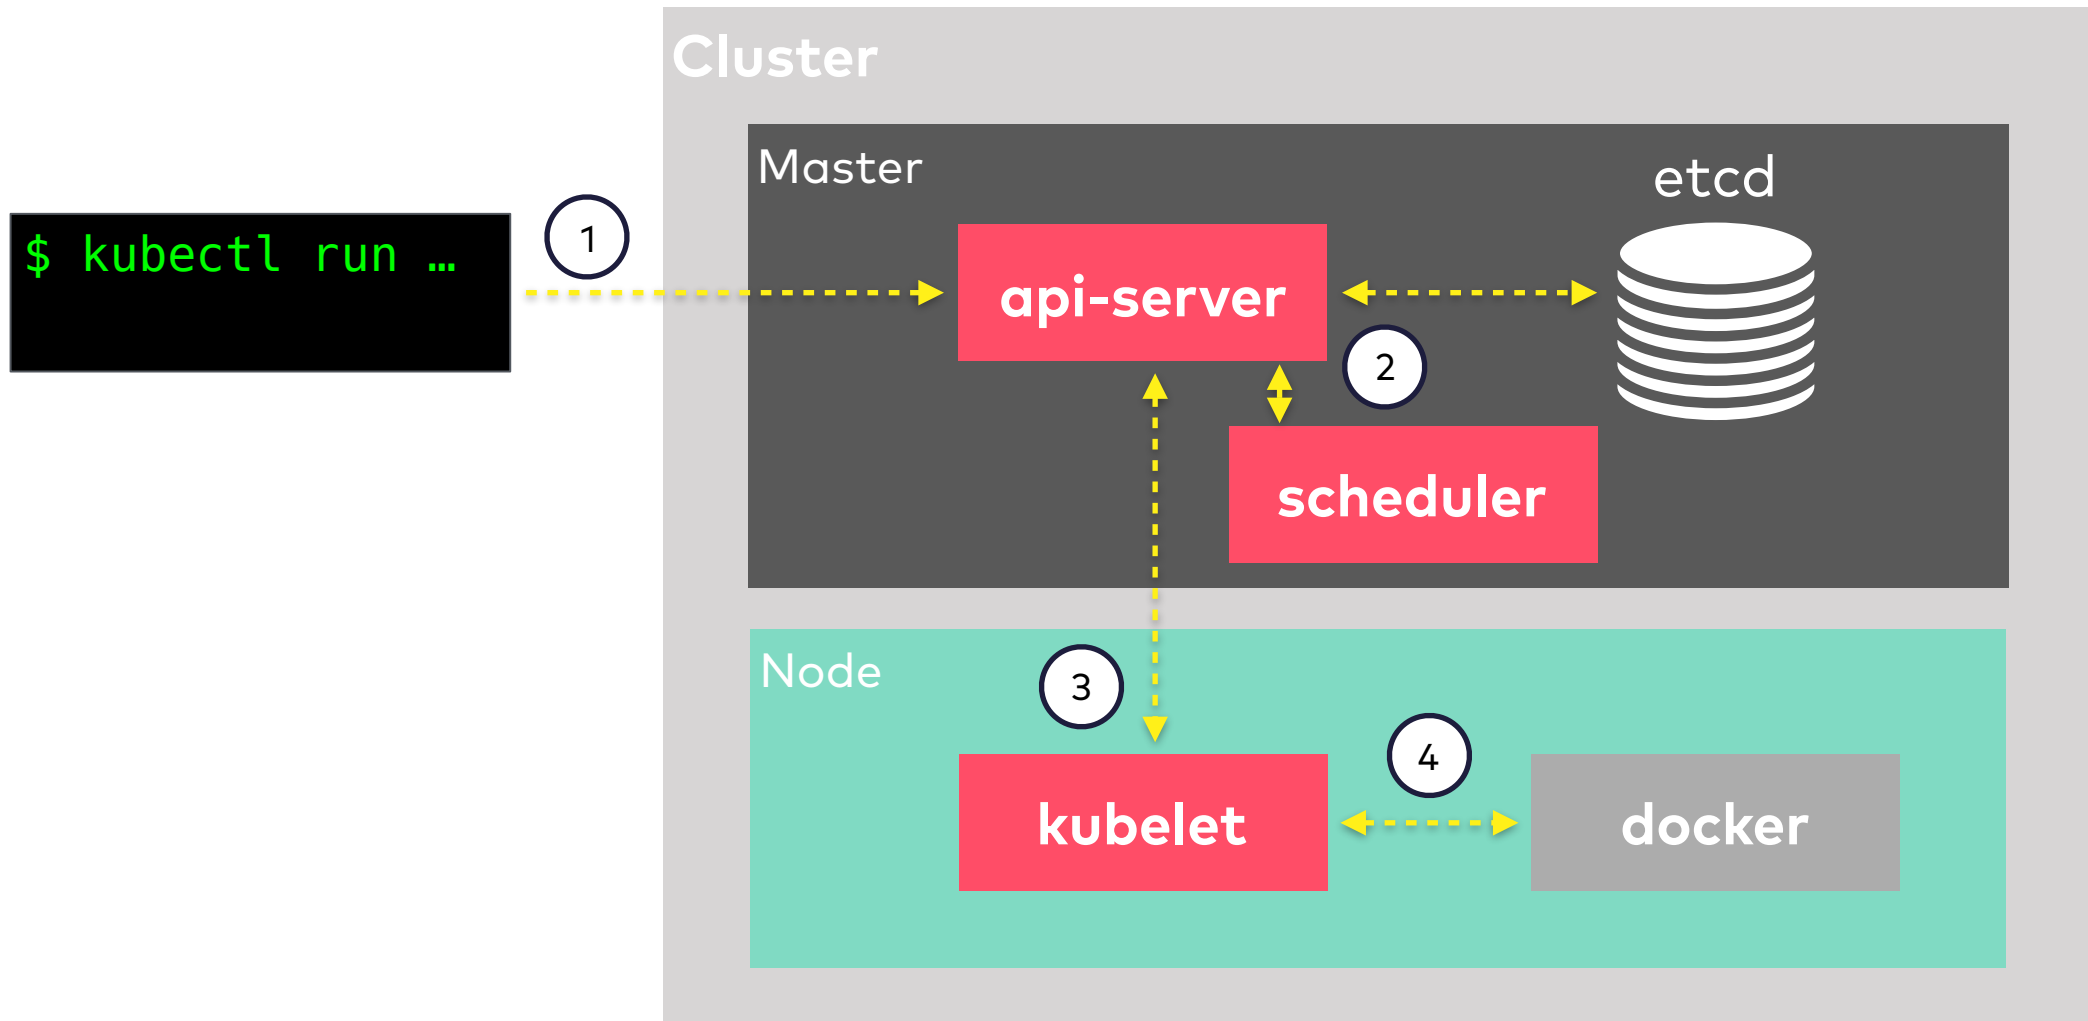


Kubelet

- Reads PodSpecs from the API
- Uses container-runtime to run Pods according to Spec

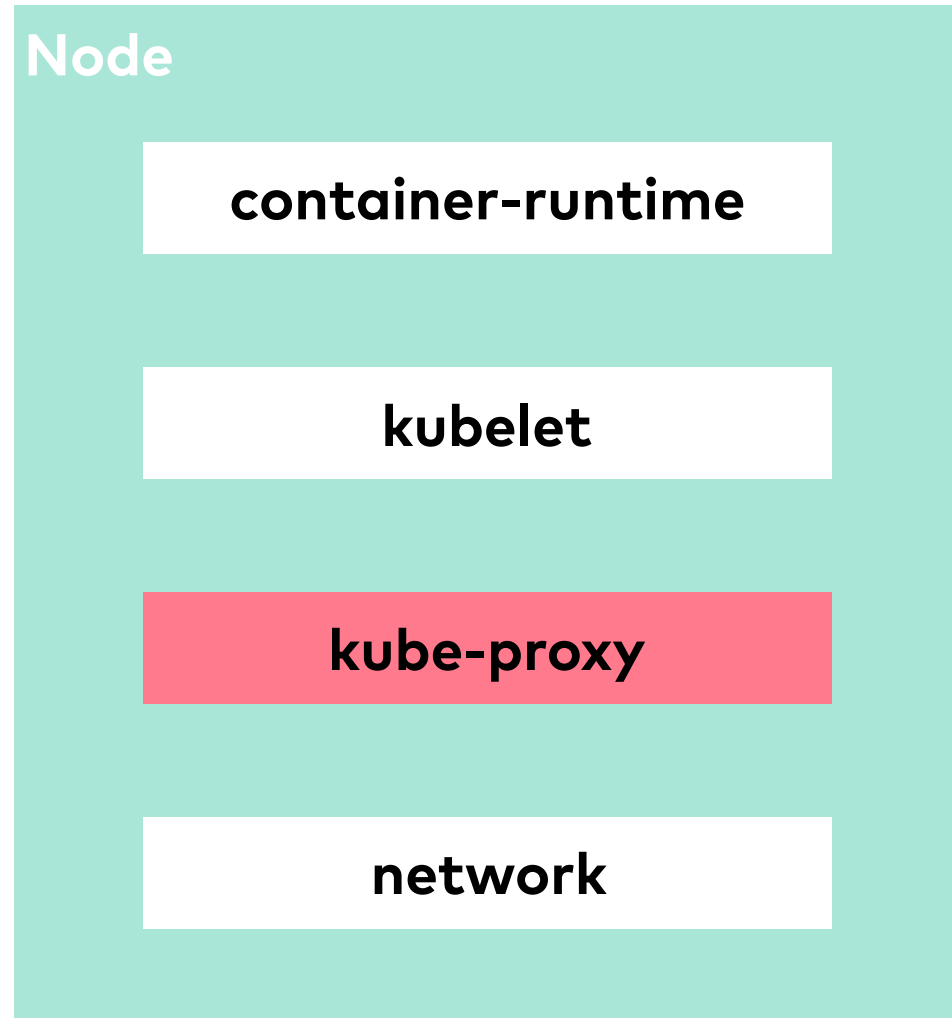


Running a Pod



Kube-Proxy

- responsible for service abstraction
- Classic proxy in usermode (old)
- New mode uses iptables to implement routing




Kube-Proxy & Services

- Service has virtual address
- Kube-proxy updates IP-Tables on Node
- Any packet with the virtual address/port combination will be changed to a node-ip and port combination
- Overlay network will then do the rest
- see IPs (172 - virtual) (10.1.x nodes)

Kube-Proxy iptables mode

Cluster



Network

- Making sure that Pods can connect across nodes
- No NAT in Cluster
- different implementations of the Container Network Interface (CNI)

Node

container-runtime

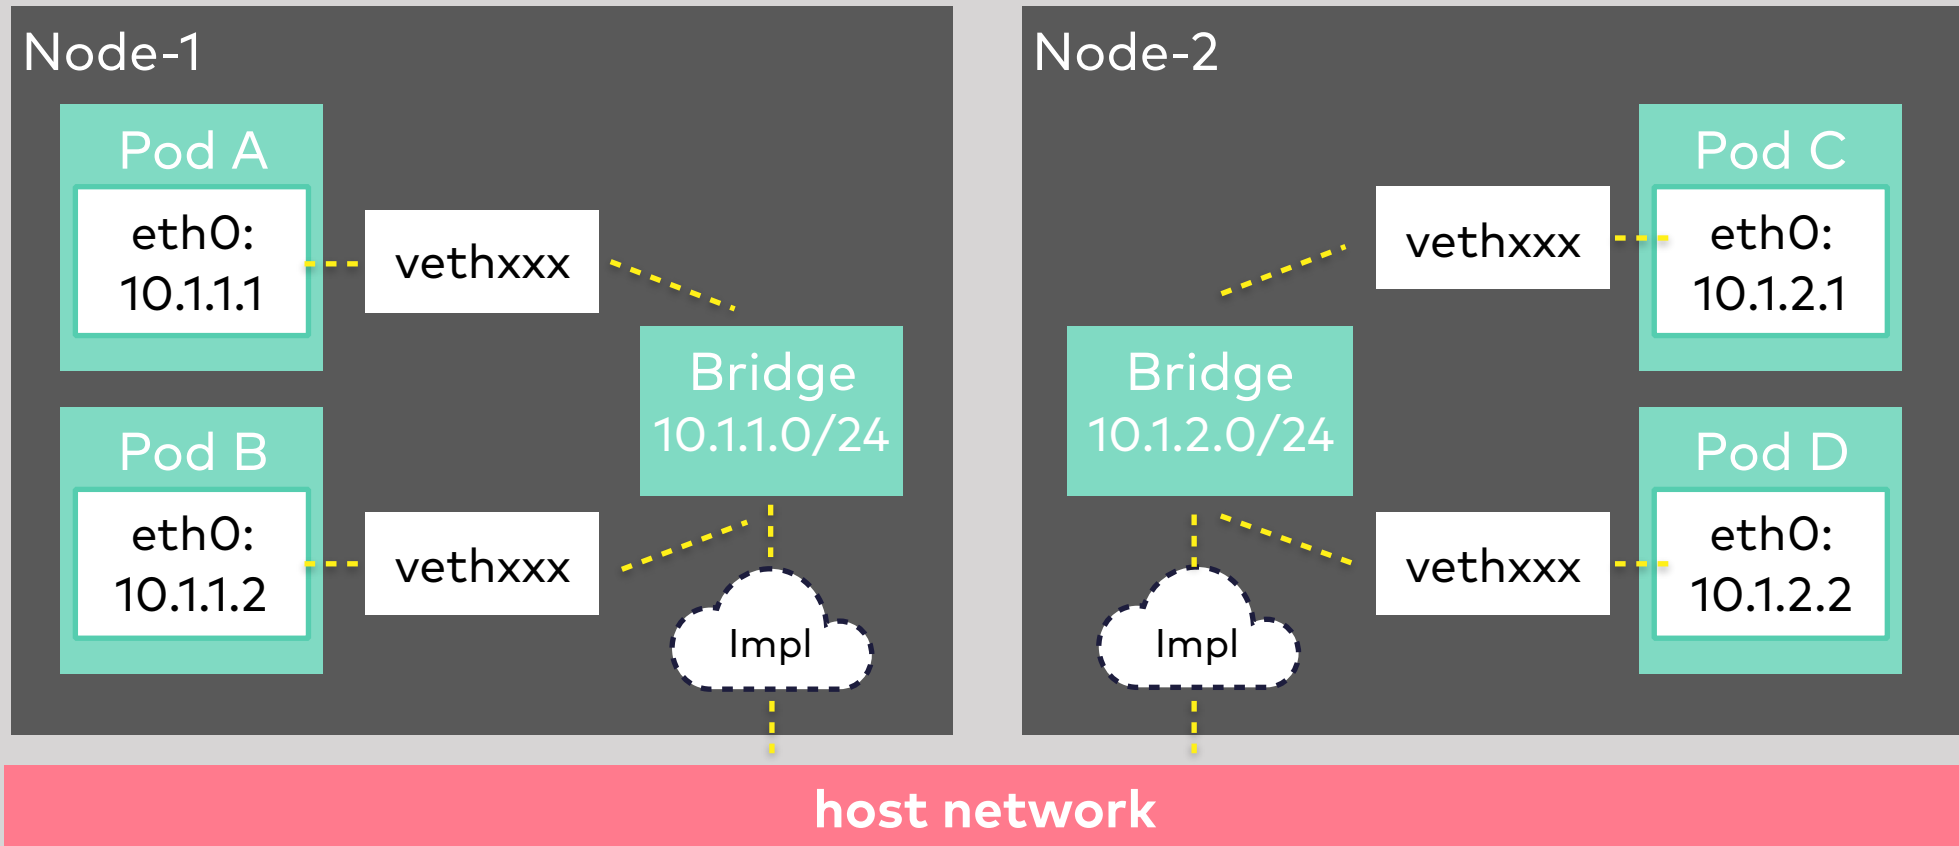
kubelet

kube-proxy

network

Network basic example

Cluster





Deploying complex applications

Helm

Helm

- Package management for Kubernetes: update, rollback, create, version, share, and publish applications
- Ready to use Kubernetes-applications (but always check the sources — like ... for real, do it)
- Handy for deployment*: persistent history and easy rollback for free
- <https://helm.sh/>

Helm cont.

- **Part of Cloud Native Computing Foundation**
- **Includes the possibility to template Kubernetes config files (e.g. for handling different clusters with the same configs)**

DEMO

[**github.com/JoergM/kubernetes_workshop_demos/helm**](https://github.com/JoergM/kubernetes_workshop_demos/helm)

Operators

Operators

- Idea to automate the knowledge of a human Operator
- Not only install automated, but **operate** automated
- The Operator itself run on Kubernetes too
- Uses Kubernetes API to do his job
- <https://coreos.com/operators/>

Advanced Features

- **Create Backups**
- **Autoscale**
- **Autoupdate installed Software**

Operator examples

- etcd
- Vault
- Prometheus
- Elasticsearch
- Kafka

Operator Framework

- Published on KubeCon 2018
- Operator SDK to create your own Operators
- Operator Lifecycle Manager - managing operators in a cluster
- Operator Metering - gathering data on operators


Service meshes

Service Meshes

- **Providing common infrastructure for microservices**
- **Features**
 - **Circuit Breaking**
 - **TLS**
 - **Authentication**
 - **Tracing**
 - **...**

Service Mesh basics

Cluster



Projects to look at

- linkerd (<https://linkerd.io/>)
- Envoy Proxy (<https://www.envoyproxy.io/>)
- Istio (<https://istio.io/>)
- Conduit (<https://conduit.io/>)

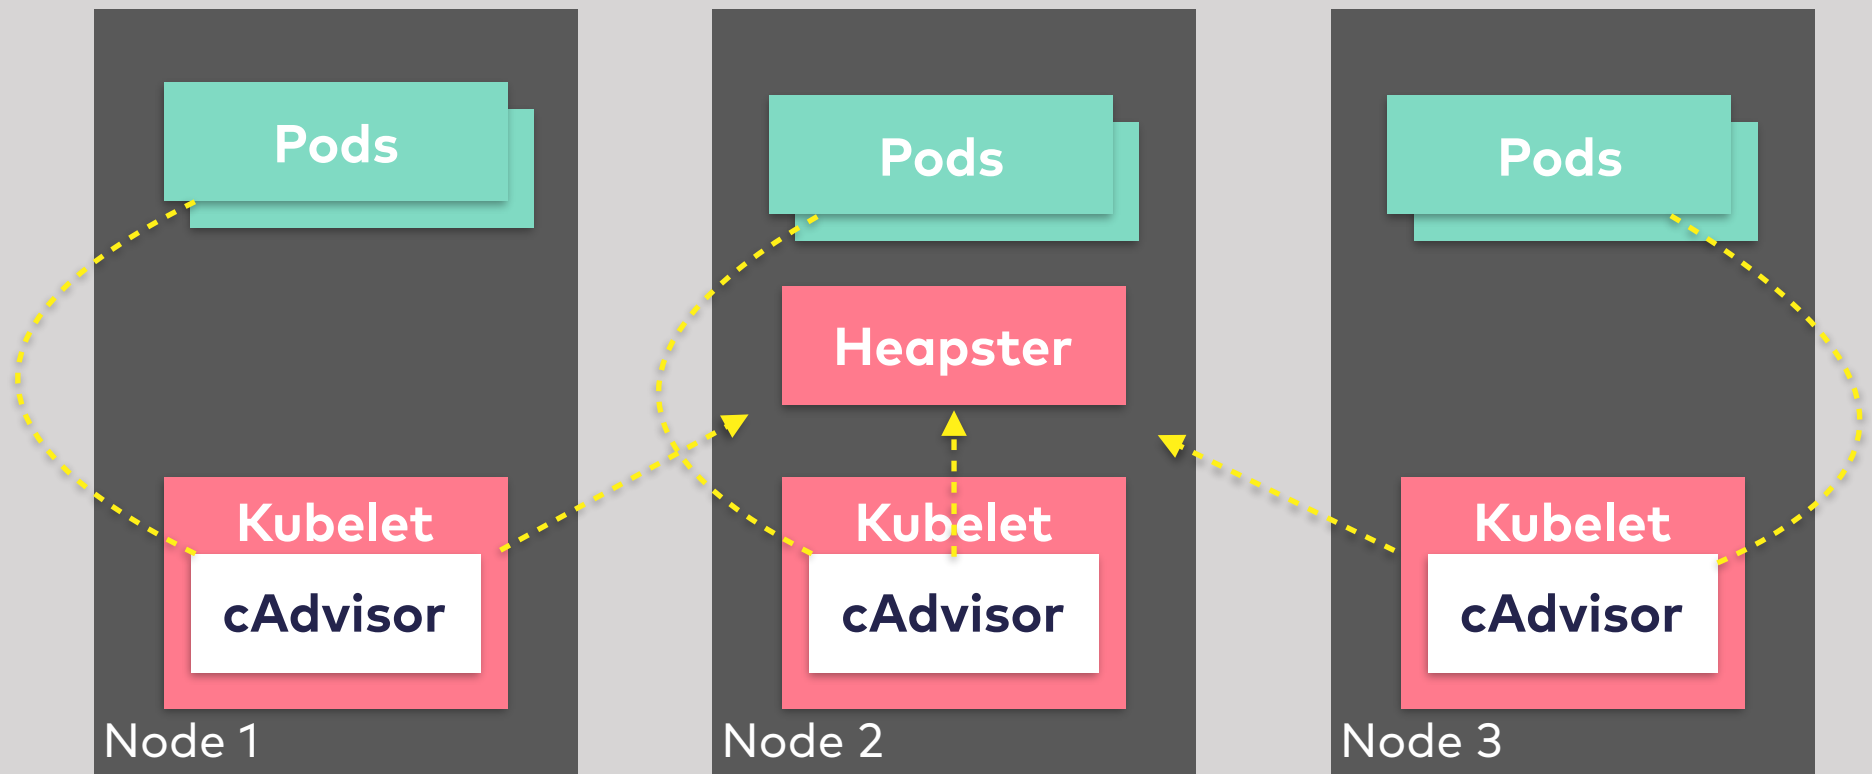


Production readiness

Monitoring

Monitoring Overview

Cluster



cAdvisor & Heapster

- cAdvisor is integrated into kubelet
 - collects performance data of containers on node
 - and on the node itself
- Heapster is running as a pod inside the cluster
 - collects data from all nodes
 - makes them available for other tools
 - Command line, dashboard, Influx, Prometheus ...

Command line

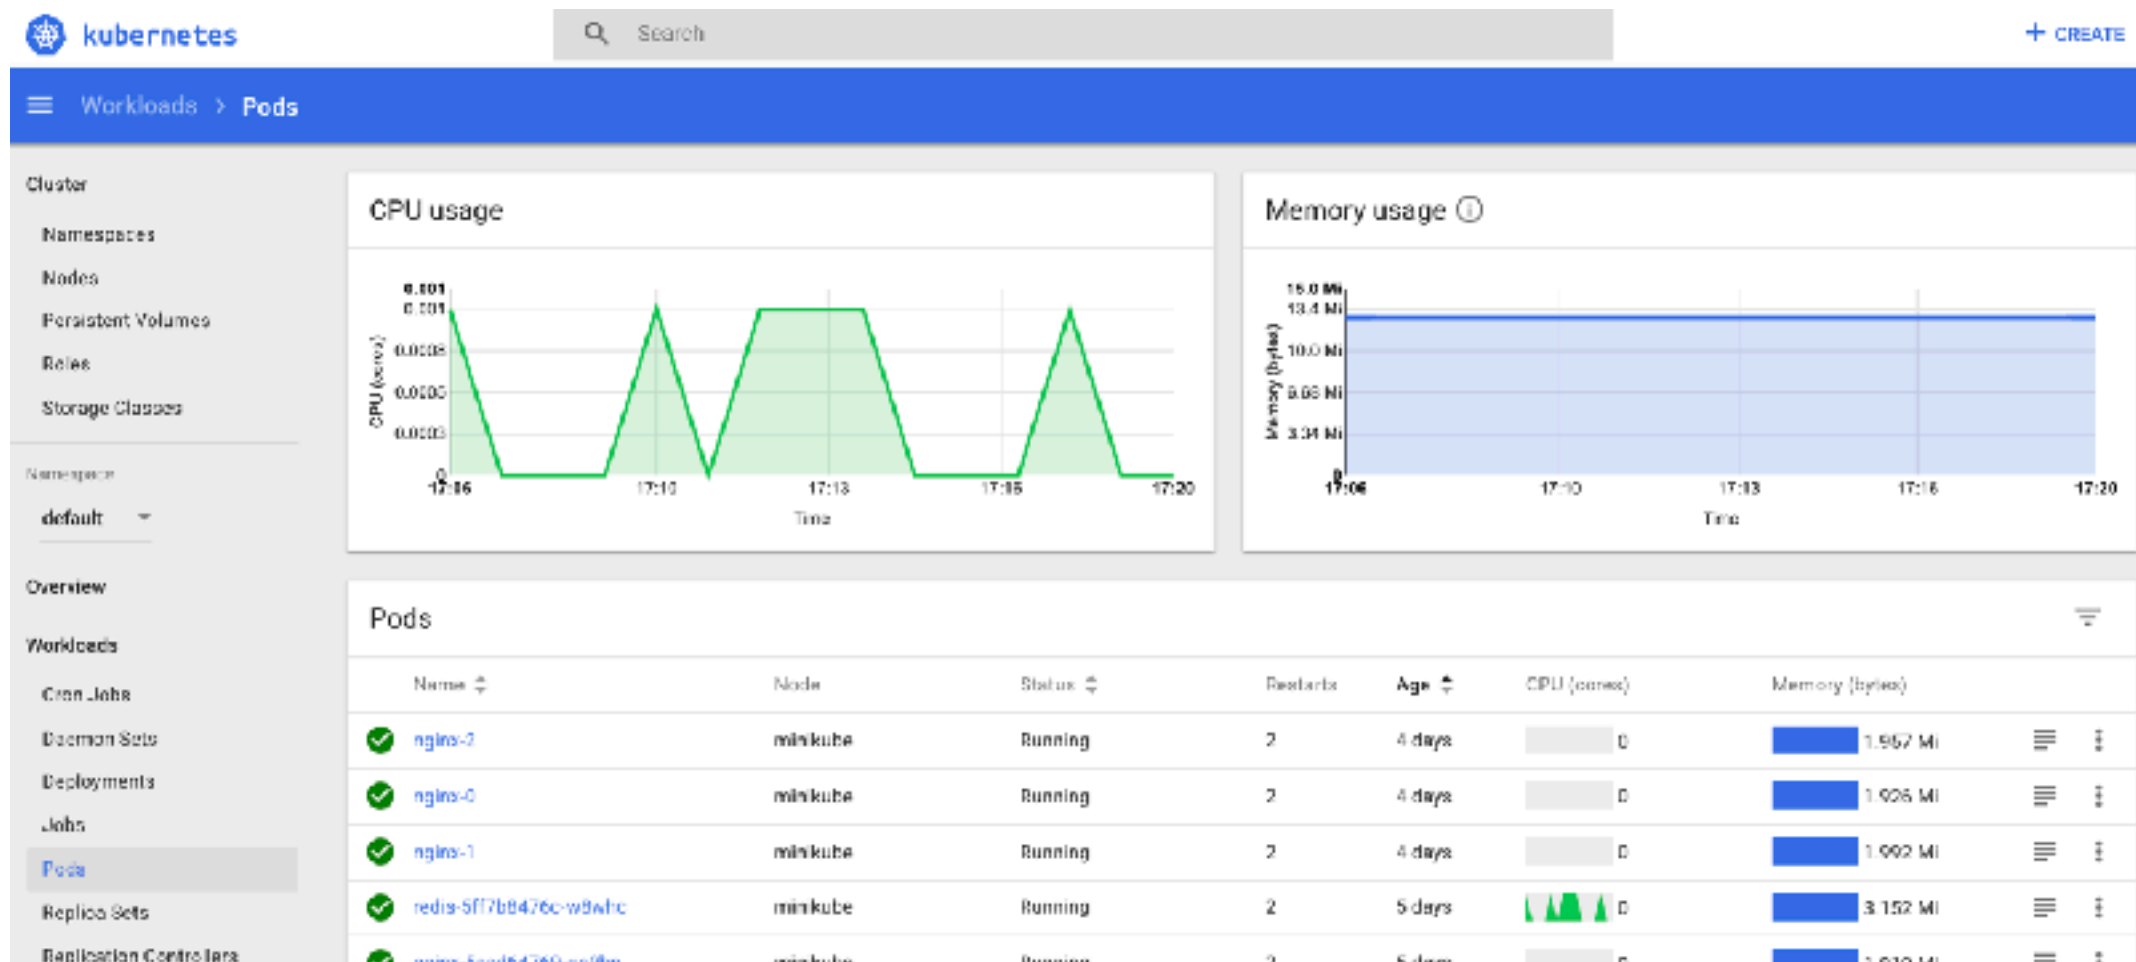
```
$ kubectl top node
```

NAME	CPU(cores)	CPU%	MEMORY(bytes)	MEMORY%
minikube	226m	11%	880Mi	22%

```
$ kubectl top pods --all-namespaces
```

NAMESPACE	NAME	CPU(cores)	MEMORY(bytes)
default	nginx-5ccd64769-gn9bn	0m	1Mi
kube-system	influxdb-grafana-rl265	2m	84Mi
kube-system	kube-addon-manager-minikube	91m	50Mi
kube-system	kube-dns-54cccfbdf8-2r526	1m	23Mi
kube-system	kubernetes-dashboard-77d8b98585-md5	3m	13Mi
...			

Dashboard



Grafana



DEMO

github.com/JoergM/kubernetes_workshop_demos/monitoring

Managing load

Ressource requests

- Requirements stated at container level
- Primarily CPU and Memory
- Scheduler uses values to find best Node

```
apiVersion: v1
kind: Pod
metadata:
  name: example-pod
spec:
  containers:
  - image: alpine
    name: foo
    resources:
      requests:
        cpu: 100m
        memory: 25Mi
```

Ressource limits

- Limits limit the resources available to a container
- If CPU exceeds limit it will be throttled
- If memory exceeds limit pod will be killed

```
apiVersion: v1
kind: Pod
metadata:
  name: example-pod
spec:
  containers:
  - image: alpine
    name: foo
    resources:
      limits:
        cpu: 100m
        memory: 25Mi
```


Limits and Requests

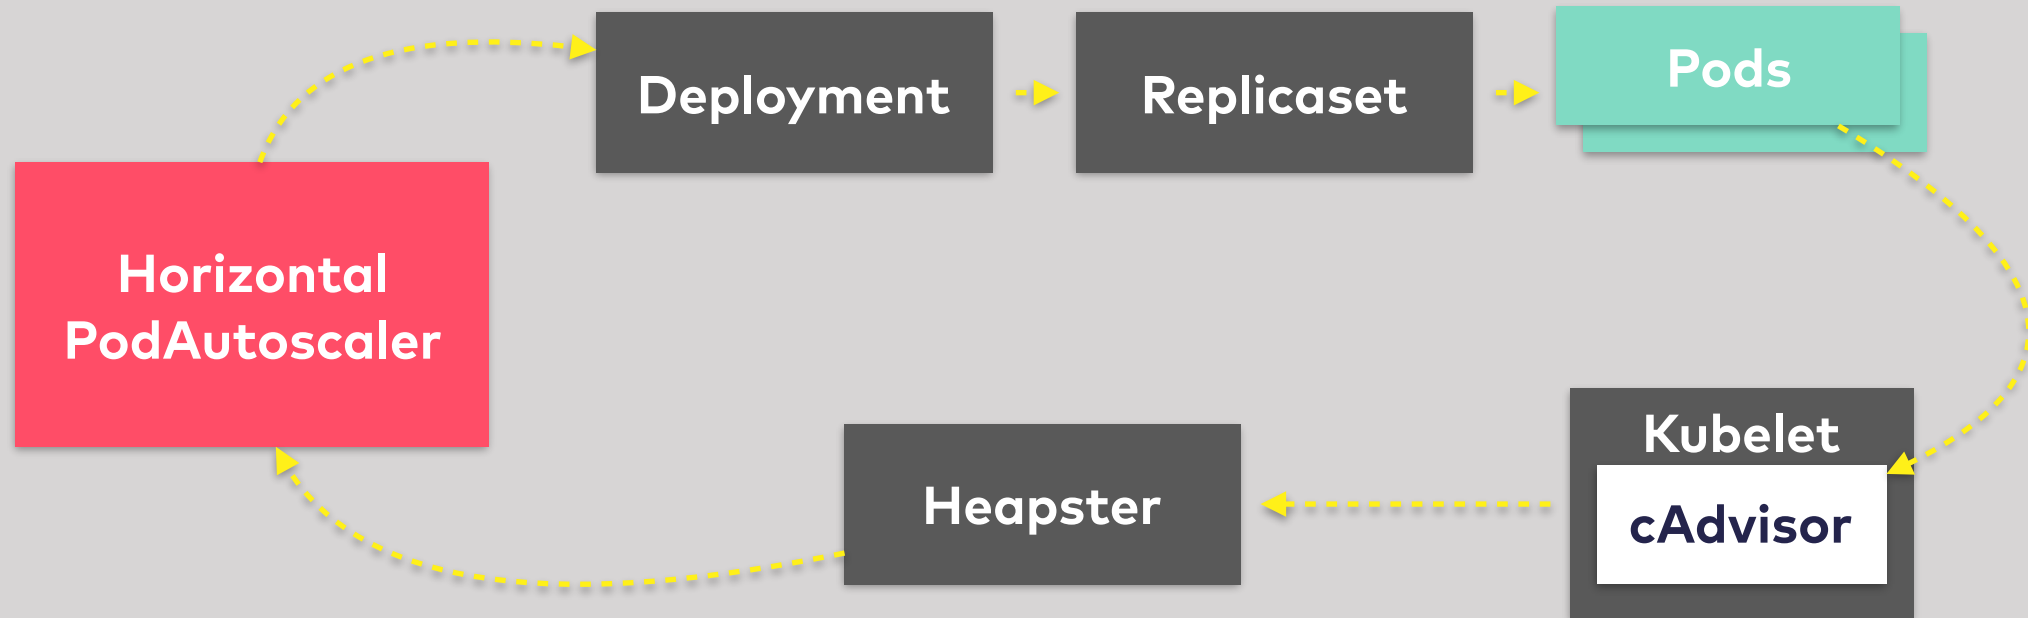
- Understand how limits and requests work
- Set them accordingly
- Be aware of resource visibility to container processes (esp. with Java applications)
 - Processes see node memory and cores

Two levels of autoscaling

- **Scaling Nodes**
 - Depending on underlying runtime environment
 - Autoscaling Groups on AWS as usual
- **Scaling Pods**
 - Cluster internal
 - Of course limited to available nodes

Scaling Overview

Cluster



Scaling Definition

- **HorizontalPodAutoscaler
API Object**
- **currently supports CPU
and custom metrics**

```
apiVersion: autoscaling/v2beta1
kind: HorizontalPodAutoscaler
metadata:
...
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: ...
  minReplicas: 1
  maxReplicas: 10
  metrics:
  - type: Resource
    resource:
      name: cpu
      targetAverageUtilization: 50
```

Security

Disclaimer

- This is only scratching at the surface
- To truly secure your cluster learn about the concepts, try them yourself, let somebody else look at it

Securing the API

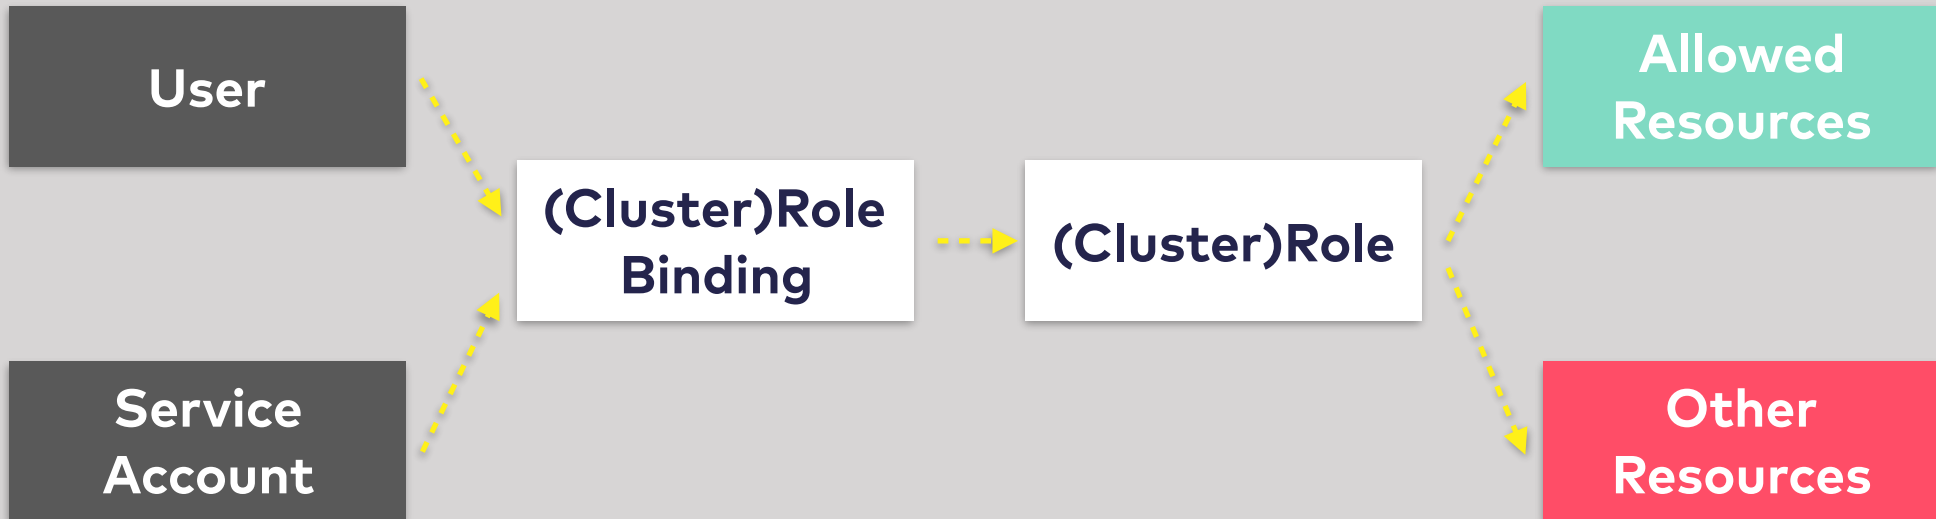
- Who is allowed to do what using the API
- Who is „Who“?
- How to identify?
- How to assign rights?

Users in Kubernetes

- **Human users**
 - **Accessing the API using e.g. kubectl**
 - **several mechanisms to identify (X.509, tokens ...)**
 - **managed externally**
- **Pods accessing the API**
 - **Pods are associated to Service Accounts**
 - **Namespace default or in Spec**

Role Based Access Control

Cluster



Pod Security Policies

- What is a Pod allowed to do?
 - User inside Pod? Is Root allowed?
 - What Kernel capabilities are allowed?
 - Read only filesystem
 - ...
- Assigned using ClusterRoles

Network Policies

- Availability depending on installed network layer
- By default every pod can be accessed (need to change)
- can isolate single pods but also namespaces

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: postgres-netpolicy
spec:
  podSelector:
    matchLabels:
      app: database
  ingress:
    - from:
      - podSelector:
          matchLabels:
            app: webserver
      ports:
        - port: 5432
```

- **architecture,
development,
devOps**
- **focus on
platform &
infrastructure**



Jörg Müller

Principal Consultant
innoQ Deutschland GmbH

joerg.mueller@innoq.com
@joergm



www.innoq.com

SERVICES

Strategy & technology consulting
Digital business models
Software architecture & development
Digital platforms & infrastructures
Knowledge transfer, coaching & trainings

FACTS

~125 employees
Privately owned
Vendor-independent

OFFICES

Monheim
Berlin
Offenbach
Munich
Zurich

CLIENTS

Finance
Telecommunications
Logistics
E-commerce
Fortune 500
SMBs
Startups